Mariya Toneva

+49 (151) 21159581
mtoneva@mpi-sws.org
http://mtoneva.com

Positions Held

Max Planck Institute for Software Systems
Tenure-track Faculty (W2) since 2022
Visiting Researcher 2021–2022

Max Planck School of Cognition
Associate Fellow since 2023

Princeton University
C.V. Starr Fellow 2021–2022
Mentors: Ken Norman, Uri Hasson

Education

Carnegie Mellon University
Ph.D. in Machine Learning and Neural Computation 2014–2021
Thesis title: Bridging Language in Machines with Language in the Brain
Advisors: Tom Mitchell, Leila Wehbe

Carnegie Mellon University
Masters of Science in Machine Learning 2018

Yale University
Bachelor of Science in Computer Science, Cognitive Science 2014

Publications in Journals and Conference Proceedings

1. Brain-tuned speech models better reflect speech processing stages in the brain 2025
   O. Moussa and **M. Toneva**
   (INTERSPEECH 2025) *Conference of the International Speech Communication Assoc.* [pdf]

2. Positional biases shift as inputs approach context window limits 2025
   B. Veseli, J. Chibane, **M. Toneva**, and A. Koller
   (COLM 2025) *Conference on Language Modeling* [pdf]

3. Improving semantic understanding in speech language models via brain-tuning 2025
   O. Moussa, D. Klakow, and **M. Toneva**
   (ICLR 2025) *International Conference on Learning Representations* [pdf]

4. Large language models can segment narrative events similarly to humans 2025
   S. Michelmann, M. Kumar, K.A. Norman, and **M. Toneva**
   Behavioral Research Methods [pdf]

5. Hints help finding and fixing bugs differently in python and text-based program representations 2025
R. Rawal, V. Padurean, S. Apel, A. Singla, and **M. Toneva**
(ICSE 2025) *International Conference on Software Engineering* [pdf]

6. Investigating the effects of fairness interventions using pointwise representational similarity 2025
C. Kolling, T. Speicher, V. Nanda, **M. Toneva**, and K.P. Gummadi
Transactions on Machine Learning Research [pdf]

7. Language models and brains align due to more than next-word prediction and word-level information 2024
G. Merlin and **M. Toneva**
(EMNLP 2024) *Empirical Methods in Natural Language Processing* [pdf]

8. Speech language models lack important brain-relevant semantics 2024
S.R. Oota, E. Çelik, F. Deniz, and **M. Toneva**
(ACL 2024) *Annual Meeting of the Association for Computational Linguistics* [pdf]

9. Perturbed examples reveal invariances shared by language models 2024
R. Rawal and **M. Toneva**
(ACL Findings 2024) *Annual Meeting of the Association for Computational Linguistics* [pdf]

10. Joint processing of linguistic properties in brains and language models 2023
S.R. Oota, M. Gupta, and **M. Toneva**
(NeurIPS 2023) *Neural Information Processing Systems* [pdf]

11. What happens during finetuning of vision Transformers: an invariance based investigation 2023
G. Merlin, V. Nanda, R. Rawal, and **M. Toneva**
(CoLLAs 2023) *Conference on Lifelong Learning Agents* [pdf]

12. Training language models for deeper understanding improves brain alignment 2023
K.L. Aw and **M. Toneva**
(ICLR 2023) *International Conference on Learning Representations* [pdf]
[top 25% notable paper (Spotlight)]

13. A Roadmap to reverse engineering real-world generalization by combining naturalistic paradigms, deep sampling, and predictive computational models 2023
P. Herholz, E. Fortier, **M. Toneva** , N. Farrugia, L. Wehbe, V. Borghesani
Neurons, Behavior, Data Science, and Theory [pdf]

14. Combining computational controls with natural text reveals aspects of meaning composition 2022
M. Toneva, T. Mitchell, and L. Wehbe
Nature Computational Science [pdf]

15. Same cause; different effects in the brain 2022
M. Toneva*, and J. Williams*, A. Bollu, C. Dann, and L. Wehbe
(CLeaR 2022) *Causal Learning and Reasoning* [pdf]

16. Single-trial MEG data can be denoised through cross-subject predictive modeling 2021
S. Ravishankar, **M. Toneva**, and L. Wehbe
Frontiers in Computational Neuroscience 2021 [pdf]

17. Modeling task effects on meaning representation in the brain via Zero-Shot MEG Prediction 2020
M. Toneva*, O. Stretcu*, B. Poczos, L. Wehbe, and T. Mitchell
(NeurIPS 2020) *Neural Information Processing Systems* [pdf]

18. Interpreting and improving natural-language processing (in machines) with natural language-processing (in the brain) 2019
M. Toneva and L. Wehbe
(NeurIPS 2019) *Neural Information Processing Systems* [pdf]

19. Inducing brain-relevant bias in natural language processing Models 2019
D. Schwartz, **M. Toneva**, and L. Wehbe
(NeurIPS 2019) *Neural Information Processing Systems* [pdf]

20. An empirical study of example forgetting during deep neural network learning 2019
M. Toneva*, A. Sordoni*, R. Tachet des Combes*, A. Trischler, Y. Bengio, and G. Gordon
(ICLR 2019) *International Conference on Learning Representations* [pdf]

21. Applying artificial vision models to human scene understanding 2015
E. M. Aminoff, **M. Toneva**, A. Shrivastava, X. Chen, I. Misra, A. Gupta, and M. J. Tarr
Frontiers in Computational Neuroscience 2015 [pdf]

22. Exploration of social grouping: effects of behavioral mimicry, appearance, and eye gaze 2014
A. Nawroj, **M. Toneva**, H. Admoni, B. Scassellati
(CogSci 2014) *Conference of the Cognitive Science Society* [with Oral presentation] [pdf]

23. The physical presence of a robot tutor increases cognitive learning gains 2012
D. Leyzberg, S. Spaulding, **M. Toneva**, and B. Scassellati
(CogSci 2012) *Conference of the Cognitive Science Society* [pdf]

24. Robot gaze does not reflexively cue human attention 2011
H. Admoni, C. Bank, J. Tan, **M. Toneva**, and B. Scassellati
(CogSci 2011) *Conference of the Cognitive Science Society* [pdf]

Preprints and Non-Proceeding Publications

- Position: episodic memory is the missing piece for long-term LLM agents 2025
M. Pink, Q. Wu, V. Vo, J. Turek, J. Mu, A. Huth, and **M. Toneva**
(arXiv 2025) [pdf]

- Assessing episodic memory in LLMs with sequence order recall tasks 2024
M. Pink, V. Vo, ... and **M. Toneva**
(arXiv 2024) [pdf]

- Vision-language integration in multimodal video transformers (partially) aligns with the brain 2023
D. Dong and **M. Toneva**
(arXiv 2023) [pdf]

- Getting aligned on representational alignment 2023
  I. Sucholutsky, L. Muttenthaler, ..., **M. Toneva**, T. Griffiths
  (arXiv 2023) [pdf]

- Interpreting multimodal video Transformers using brain recordings 2023
  D. Dong and **M. Toneva**
  (ICLR 2023 Workshop on Multimodal Representation Learning: Perks and Pitfalls) [pdf]

- Memory for long narratives 2022
  M. Toneva, V. Vo, J. Turek, S. Jain, S. Michelmann, M. Capotă, A. Huth, U. Hasson, and K. Norman
  (CEMS 2022) *Context and Episodic Memory Symposium*

- The Courtois Neuromod project: a deep, multi-domain fMRI dataset to build individual brain models 2022
  J. Boyle*, B. Pinsard*, ... **M. Toneva**, ... and P. Bellec
  (HBM 2022) *Human Brain Mapping* [with Oral presentation]

- Does injecting linguistic structure into language models lead to better alignment with brain recordings? 2021
  M. Abdou, A.V. González, **M. Toneva**, D. Hershcovich, and A. Søgaard
  (arXiv 2021) [pdf]

- Investigating different alignment methods between natural and artificial neural networks for language processing 2020
  A. Bollu, **M. Toneva**, and L. Wehbe
  (SNL 2020) *Society for the Neurobiology of Language*

- Investigating task effects on brain activity during stimulus presentation in MEG 2019
  M. Toneva*, O.Stretcu*, B. Poczos, and T. Mitchell
  (HBM 2019) *Human Brain Mapping*

- Word length processing in left lateraloccipital through region-to-region connectivity: an MEG Study 2018
  M. Toneva, and T. Mitchell
  (HBM 2018) *Human Brain Mapping*

- MEG representational similarity analysis implicates hierarchical integration in sentence processing 2018
  N. Rafidi*, D. Schwartz*, **M. Toneva***, S. Jat, and T. Mitchell
  (HBM 2018) *Human Brain Mapping*

- Scene-space encoding within the functional scene-selective network 2015
  E. M. Aminoff, **M. Toneva**, A. Gupta, and M. J. Tarr
  (VSS 2015) *Vision Sciences Society*

- Towards a model for mid-level feature representation of scenes 2014
  M. Toneva, E. M. Aminoff, A. Gupta, and M. Tarr
  (VSS 2014) *Vision Sciences Society*

Grants and Fellowships

DFG Graduate School RTG 2853
Neuroexplicit Models of Language, Vision & Action 2023-2028

DFG Research Unit 5368 KI-FOR
Abstract Representations in Neural Architectures 2023-2027
Project: Bridging Levels of Abstraction in Brains and Natural Language Processing Machines

National Institutes of Health T32 Training Grant
Funded one year of postdoctoral research in Quantitative and Computational Neuroscience 2021-2022

C.V. Starr Fellowship
Funded one year of postdoctoral research in computational neuroscience at Princeton University 2021-2022

National Science Foundation Graduate Research Fellowship
Funded three years of interdisciplinary graduate research in machine learning and neuroscience 2016–2019

Grace Hopper Celebration Scholarship
Funded attendance at the 2014 Grace Hopper Celebration of Women in Computing 2014

Mellon Forum Undergraduate Research Grant
Funded submission and attendance at the 2014 Vision Sciences Society conference 2014

Robin Berlin Fellowship
Funded neural modeling research at Laboratory of Computational Neuroscience, EPFL 2013

Awards and Distinctions

Japanese-American-German Frontiers of Science (JAGFOS) Symposium Invitee
Alexander von Humboldt Foundation 2023
Selected to represent Germany as one of 24 researchers from all disciplines of science and engineering

Appointed as Member
ELLIS (European Laboratory for Learning and Intelligent Systems) 2023

Appointed as Adjunct Fellow
Max Planck School of Cognition 2023

Ph.D. Dissertation Award, Honorable Mention
Society for the Neurobiology of Language 2021

Machine Learning Student Leadership Award
Awarded for exemplary efforts and their significant impact on life in the Machine Learning Department 2020

Top Reviewer
NeurIPS 2022, ICLR 2022, ICML 2022, NeurIPS 2018 2018-2022

Citadel Datathon Runner-up
Analyzed a genomics dataset to predict age-related differences in disease-related gene expression 2017

Machine Learning Teaching Assistant Award
Awarded for outstanding performance as a TA in 10-725 Convex Optimization 2017

BrainHub Neurohackathon Winner
Reduced need for human supervision by classifying diffusion MRI tracks into anatomical bundles 2016

Invited and Contributed Talks

Conferences and Workshops

- **Keynote Speaker**, ELLIS x UniReps Speaker Series 2025
- Panel member, ICLR Workshop on "New Frontiers in Associative Memory" 2025
- **Keynote Speaker**, CALCULUS Symposium, KU Leuven 2024
- Invited Speaker, Brain Prize Webinar "LLMs and Human Language Processing" 2024
- **Keynote Speaker**, AAAI Workshop "AI for Brain Encoding and Decoding" 2024
- Invited Speaker, ICLR Workshop "Representational Alignment" 2024
- Invited Speaker, CogSci Symposium "Is DL the Answer for Understanding Cog. Dynamics" 2024
- Invited Speaker, German Academy of Science Symposium "Brain Science and LLMs" 2024
- Invited Speaker, NeuroAI Symposium 2024
- **Keynote Speaker**, Neuro-AI Talks (NEAT) 2023
- Invited Speaker, Workshop on Philosophy of Science Meets ML, University of Tübingen 2023
- Invited Symposium Speaker, European Society for Philosophy and Psychology Conference 2023
- Invited Speaker, Neuro-AI Educational Session, Human Brain Mapping Conference 2023
- Invited Speaker, Workshop on Code, Brains, and LLMs, Saarland University 2023
- Invited Speaker, Workshop on ML, Abstract Thought, and the Expanding Reach of AI 2022
- Contributed Talk, Neuromatch Conference 2020
- Invited Speaker, CCN Workshop on Nonlinear Models for Scientific Discovery 2020
- Contributed Talk, NeurIPS Workshop Women in Machine Learning 2014
- Contributed Talk, CogSci Oral Presentation for Accepted Paper 2014

University and Industry Seminars

- Institute Colloquium, Gatsby Computational Neuroscience Unit 2025
- Cohere AI Social 2025
- Institute Colloquium, Ernst Strüngmann Institut 2024
- Math Machine Learning Seminar, Max Planck Institute for Mathematics and UCLA 2024
- Institute Colloquium, MPI for Human Cognitive and Brain Sciences 2024
- Campus Lecture, Saarbrücken Informatics Campus 2023
- Cog. Comp. Neuro. Colloquium, MPI for Human Cognitive and Brain Sciences 2023
- Psychology Department Colloquium, University of Saarland 2023
- Institute Colloquium, Institute for Basic Science, South Korea 2023
- Institute Colloquium, CIMeC, University of Trento 2022
- Distinguished Speakers in Language Science Colloquium, Saarland University 2022
- Department of Statistics and Data Science, Yale University 2021
- Faculty of Computing and Data Science, Boston University 2021
- Department of Computer Science, University of Southern California 2021
- Department of Computer Science, University of Utah 2021
- Institute Colloquium, Max Planck Institute for Software Systems 2021

- Department of Computer Science, Aarhus University 2021
- Institute Colloquium, Toyota Technological Institute at Chicago 2021
- Institute Colloquium, IST Austria 2021
- Department of Statistics, University of Chicago 2021
- Department of Computer Science, University of Notre Dame 2021
- Department of Computer Science, University of Liverpool 2020
- SFB-TRR 161 Lecture Series, U. of Stuttgart, U. of Konstanz, Ulm University, and LMU 2020

Summer Schools and Group Meetings

- Heinrich Heine University Düsseldorf, Host: Milica Gasic 2024
- Goethe University Frankfurt, Host: Christian Fiebach 2024
- Bernstein Center for Computational Neuroscience Retreat 2024
- Max Planck School of Cognition Academy 2023
- Summer School in Philosophy and Computer Science, University of Bayreuth 2023
- Department of Computer Vision and Machine Learning, MPI for Informatics 2023
- TALEP group, Aix-Marseille University 2022
- CMMRS Summer School, MPI for Software Systems 2022
- IMPRS NeuroCom Summer School at the MPI for Human Cognitive and Brain Sciences 2022
- MIT, Host: Evelina Fedorenko 2022
- MIT, Host: Roger Levy 2021
- Computational Neuroscience Symposium, CMU 2021
- Courtois NeuroMod Group, University of Montreal 2021
- Princeton Neuroscience Institute, Hosts: Ken Norman and Uri Hasson 2021
- brAIn seminar, CMU 2020
- UT Austin, Host: Alexander Huth 2020

Mentorship and Supervision

Omer Moussa
PhD at University of Saarland 2024-

Camila Kolling
PhD at University of Saarland 2024-
co-advised by Krishna Gummadi

Mathis Pink
PhD at University of Saarland 2024-
co-advised by Isabel Valera

Emin Çelik
Postdoctoral Researcher 2023-

Blerta Veseli
PhD at University of Saarland 2023-
co-advised by Alexander Koller

Cameron Braunstein
PhD at University of Saarland 2023-

co-advised by Eddy Ilg

Gabriele Merlin
PhD at CS@MaxPlanck Graduate Program 2022-

Shashwat Saxena
Research Intern; now Masters student at CMU 2025

Michela Proietti
Research Intern 2025

Alan Sun
Research Intern; now Masters student at CMU 2024

Khai Loong Aw
Research Intern; now PhD student at Stanford University 2022

Ruchit Rawal
Research Intern; now PhD student at University of Maryland 2022-2024

Subba Reddy Oota
Research Intern; now postdoc at TU Berlin 2022

Tianai (Dota) Dong
Masters Student; now PhD student at the MPI for Psycholinguistics 2021-2022

Teaching

Bridging Language in Machines and Language in the Brain, University of Saarland
Instructor 2023
Seminar course

3370 Mathematical Neuroscience, University of Pittsburgh
Teaching Assistant 2018

10-725 Convex Optimization, CMU
Teaching Assistant 2016
Awarded Machine Learning TA award

Machine Learning for Neuroscience, Multimodal Neuroimaging Training Program
Instructor 2016
Created curriculum and instructed 4-week course; video recordings can be found on personal webpage

Service

Organizer

Cajal Neuroscience and AI Summer School 2025
Champalimaud Centre for the Unknown

Deep Learning for Brain Encoding and Decoding 2023
Tutorial, IJCAI

Memory in Artificial and Real Intelligence 2022
Workshop, NeurIPS

Deep Learning for Brain Encoding and Decoding 2022
Tutorial, Cognitive Sciences Society

What can NLP systems teach us about language in the brain? 2021
Symposium, Society for the Neurobiology of Language

How can findings about the brain improve AI systems? 2021
Workshop, ICLR

Program Committee

Program Chair: *CogSci* 2024
Senior Program Committee Member: *ACL Rolling Review* (2023–present), *NeurIPS* (2024–present), *ICML* (2025–present), *CCN* (2025-present), *CCN* Technical Program Committee (2022–2024)
Program Committee Member: ML: *NeurIPS* 2016-2023(Top 30% Reviewer in 2018); *ICML* 2019-2023 (Top 10% Reviewer in 2022); *AAAI* 2020-2021, *CoLLAs* 2022, *ICLR* 2022-2024 (Highlighted Reviewer in 2023); **NLP:** *ACL* 2019-2021; *NAACL* 2019-2021; *EMNLP* 2020-2021; *CoNLL* 2020-2021; *AACL-IJCNLP* 2020; *EACL* 2021
Journal Reviewer: *TMLR*, *Nature Human Behavior*, *Nature Communications*; *Communications Biology*; *TICS*, *Communications of the ACM*, *Frontiers in Computational Neuroscience*

Thesis Committee Member

Viktor Kewenig (UCL, 2025); Aylin Kallmayer (Goethe University Frankfurt, 2025); Till Speicher (MPI Software Systems, 2025); Carina Kauf (MIT, 2024); RJ Antonello (UT Austin, 2024); Julien Dirani (NYU, 2024); Bernhard Schäfl (Johannes Kepler University, 2024); Damián Pascual (ETH Zurich, 2022)

ML@CMU Blog

Chief Editor and Co-founder 2018–2020
Oversaw more than 30 research posts featuring recent ML research across 6 departments in the School of Computer Science as well as other CMU schools and departments, and more than 10 educational posts

University Leadership Student Advisory Council

Member 2015–2017
Advising senior leadership at Carnegie Mellon University on the strategic priorities of the university

Graduate Student Assembly

Representative for the Program of Neural Computation 2015–2018
Advocating for the needs of graduate students

Yale Review of Undergraduate Research in Psychology

Chief Editor 2013–2014
Reviewed 50 submissions from 31 universities, and edited 9 submissions for publication

Industry Internships

Microsoft Research, Montreal

Research Intern 2018
Investigated the learning dynamics of neural networks as they train on single classification tasks, finding that certain examples are forgotten with high frequency, and some not at all, and that, based on these forgetting dynamics, a significant fraction of examples can be omitted from the training data set while still maintaining state-of-the-art generalization performance

Cognitive Computing Center, Thomson Reuters

Research Intern 2017
Investigated the use of a recurrent neural network encoder for unsupervised word-order sensitive hashing as a step towards improving ranking results

Research Visits

Carnegie Mellon University
Research Assistant; Advisor: Michael Tarr 2013–2014
Investigated mid-level scene representation in humans using computer vision techniques

École Polytechnique Fédérale de Lausanne (EPFL)
Summer Intern; Advisor: Wulfram Gerstner 2013
Worked towards improving the state-of-the-art calcium-based model of spike-timing dependent plasticity

Massachusetts Institute of Technology
Technical Trainee; Advisor: John Gabrieli 2012
Examined links between working memory capacity and various brain metrics through the analysis of resting state functional connectivity fMRI data

Personal

Languages	Bulgarian (Native), English (Fluent), German (Intermediate)
Citizenship	United States, Bulgaria
Github profile	http://github.com/mtoneva
Google Scholar profile	https://scholar.google.com/citations?user=a61sk-4AAAAJ