

Please pick up an anonymous  
evaluation sheet

# Homework 3 Questions

## 1 K-means

Run examples 1 through 3 in the attached code. Each example displays a different result of running the k-means clustering algorithm on the iris data set. Which result do you think is the best, and why do the other runs of k-means fail?

## 2 PCA

How many principal components (PCs) could possibly be in the iris data set?

Run example 4 in the attached code. How much of the variance in the data set is explained by the first PC?

Run example 5 in the attached code. Does including another component improve the results? Why or why not?

Question 1: When doing PCA, using a higher number of PCs to create the final reduced representation of the data would \_\_\_\_\_ overfitting.

- a. increase
- b. decrease

Question 2: CCA can be used to find the linear transformations of more than 2 data sets that are most correlated.

- a. true
- b. false

Question 3: You want to perform k-means 2 separate times. You have asked for 2 clusters. The data also has labels, though note that k-means is still unsupervised (i.e. does not make use of the labels). After completing the first k-means with k-means++ initialization, you observe that all data from class 1 is in cluster 1. Now, you run k-means a second time, but you specify the means of the clusters to be the means of the clusters obtained after the first run of k-means. Do you expect that all data points from class 1 will again be in cluster 1?

- a. yes
- b. no

Question 4: Now, instead of running a second k-means with the cluster means initialized to be the cluster means from the first run, you run a second k-means with k-means++ initialization. Do you expect that all data points from class 1 will again be in cluster 1?

- a. yes
- b. no

# Advanced topics: latent variable models, deep learning, and reinforcement learning

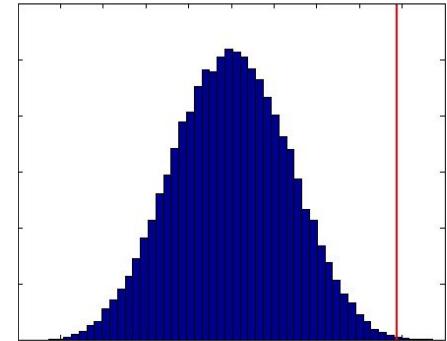
06/24/2016

Mariya Toneva

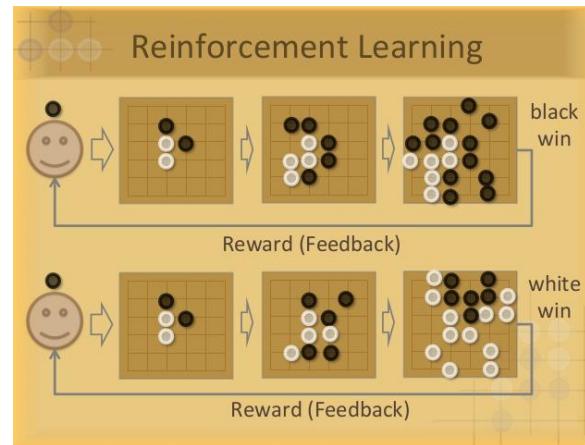
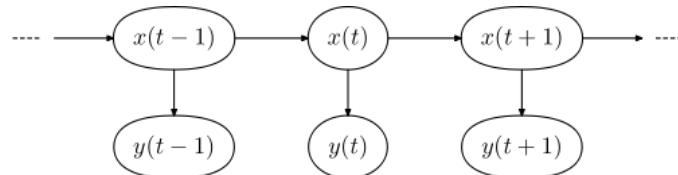
[mariya@cmu.edu](mailto:mariya@cmu.edu)

Some figures derived from slides  
by Tom Mitchell, Aarti Singh,  
Barnabás Póczos

# How can ML help neuroscientists?



- ❑ Evaluate results
  - ❑ cross validation (how generalizable are our results?)
  - ❑ nearly assumption-free significance testing (are the results different from chance?)
- ❑ Complex data-driven hypotheses of brain processing
  - ❑ advanced topics: latent variable models, reinforcement learning, deep learning



Look ahead to end of lecture: understanding the most recent AI feat of beating Go world champion



Mnih et al., 2013



Silver et al., 2016 Nature

# Today: advanced topics!

- ❑ Latent variable models
  - ❑ Hidden Markov models
- ❑ Reinforcement learning
- ❑ Deep learning
  - ❑ Neural networks
    - ❑ Recurrent NNs
      - ❑ Long short-term memory networks
    - ❑ Deep belief networks
    - ❑ Convolutional neural networks
- ❑ AlphaGo: RL + deep learning!

# Importance of sequential data

- ❑ Moving from treating all observations as independent from each other, to dependent

# Importance of sequential data

- ❑ Moving from treating all observations as independent from each other, to dependent
- ❑ Important, when considering sequential data, like we have in neuroscience:

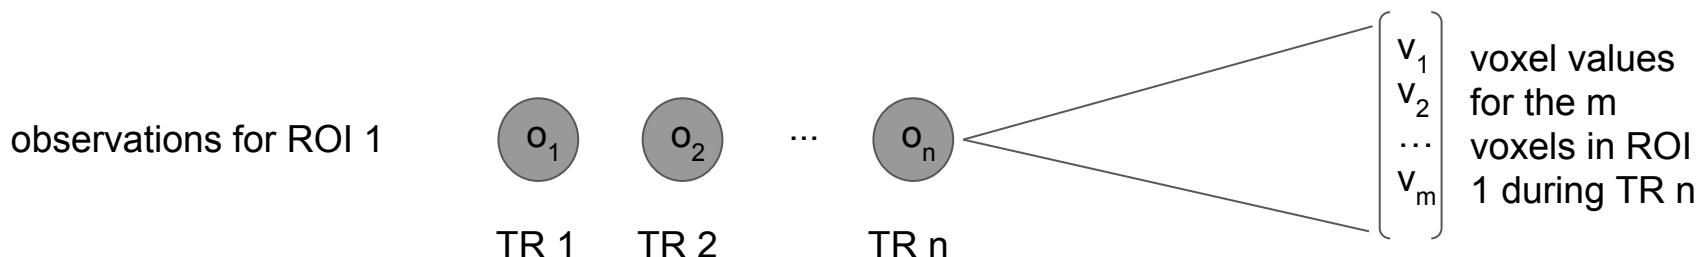
# Importance of sequential data

- ❑ Moving from treating all observations as independent from each other, to dependent
- ❑ Important, when considering sequential data, like we have in neuroscience:
  - ❑ Example: observation 1 = fMRI voxel activities in ROI 1 for TR 1
  - ❑ observation 2 = fMRI voxel activities in ROI 1 for TR 2
  - ❑ ...
  - ❑ observation N = fMRI voxel activities in ROI 1 for TR N



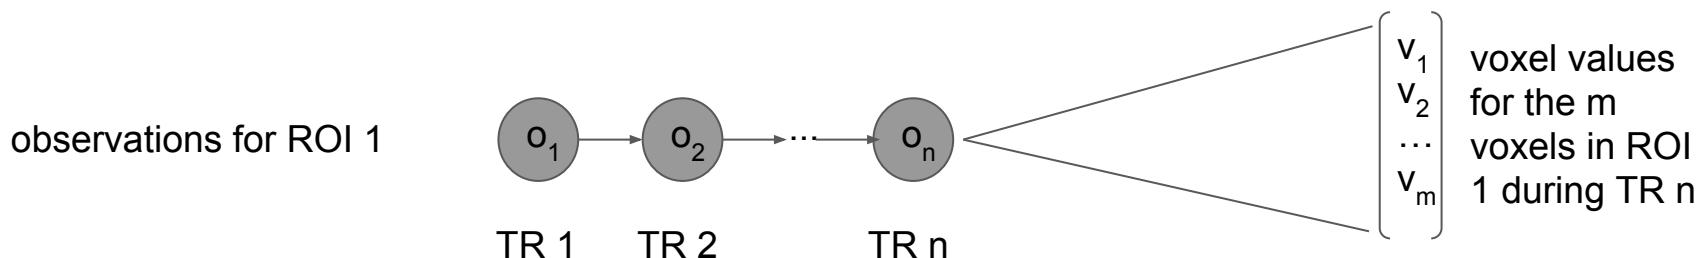
# Importance of sequential data

- ❑ Moving from treating all observations as independent from each other, to dependent
- ❑ Important, when considering sequential data, like we have in neuroscience:
  - ❑ Example: observation 1 = fMRI voxel activities in ROI 1 for TR 1
  - ❑ observation 2 = fMRI voxel activities in ROI 1 for TR 2
  - ❑ ...
  - ❑ observation N = fMRI voxel activities in ROI 1 for TR N



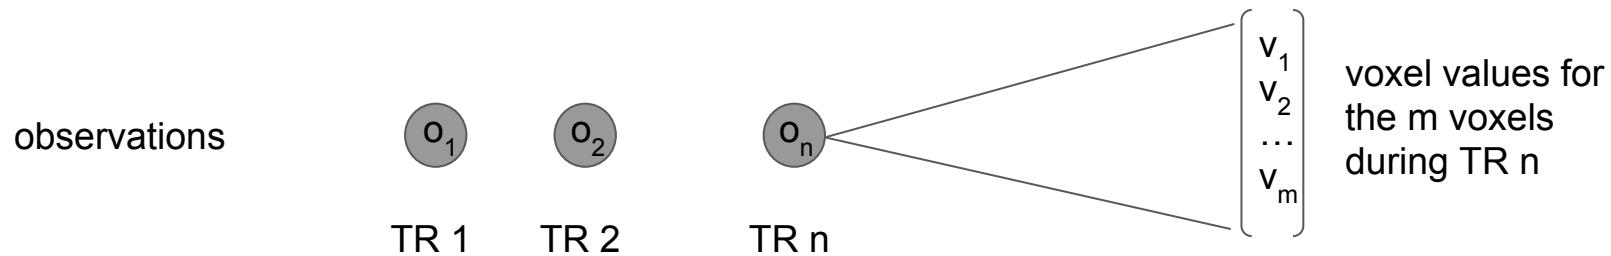
# Importance of sequential data

- ❑ Moving from treating all observations as independent from each other, to dependent
- ❑ Important, when considering sequential data, like we have in neuroscience:
  - ❑ Example: observation 1 = fMRI voxel activities in ROI 1 for TR 1
  - ❑ observation 2 = fMRI voxel activities in ROI 1 for TR 2
  - ❑ ...
  - ❑ observation N = fMRI voxel activities in ROI 1 for TR N



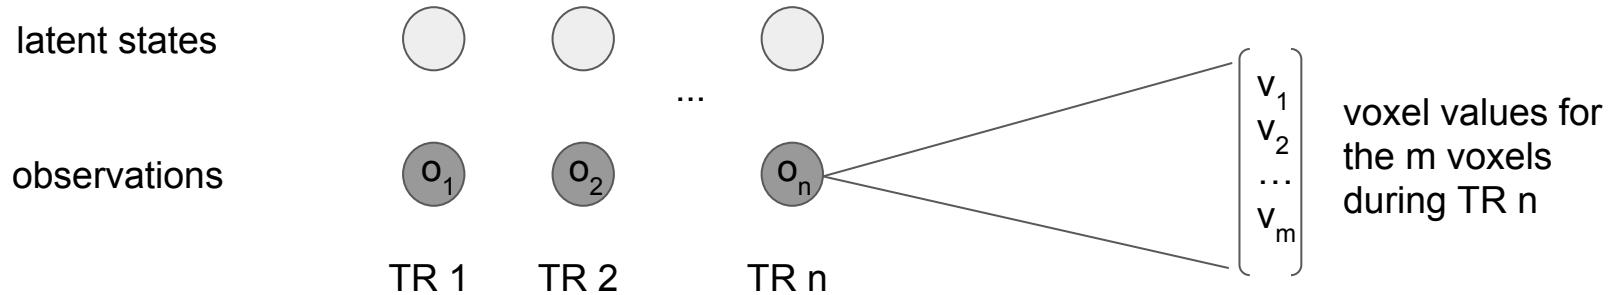
# What if there are latent processes?

- Dependence in observations may be due to dependence in latent states



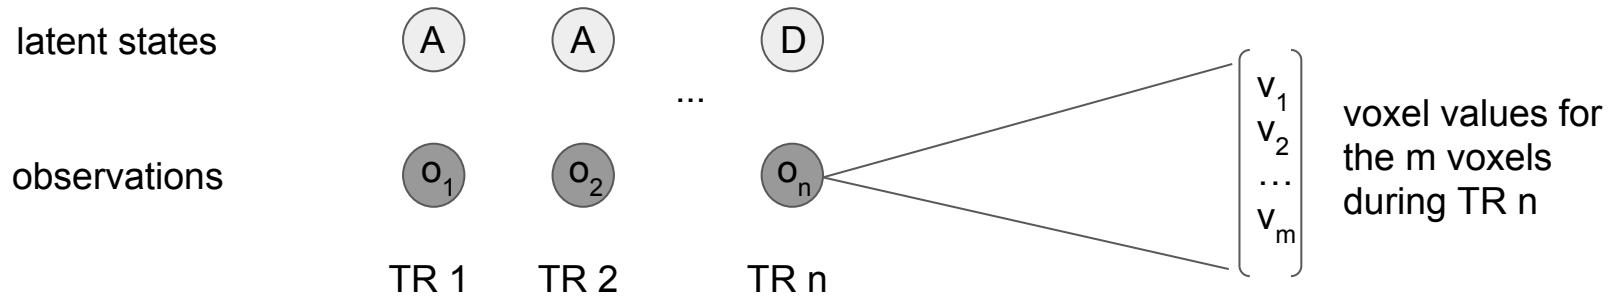
# What if there are latent processes?

- Dependence in observations may be due to dependence in latent states
  - Example:
    - underlying processes while subject is at rest in an fMRI



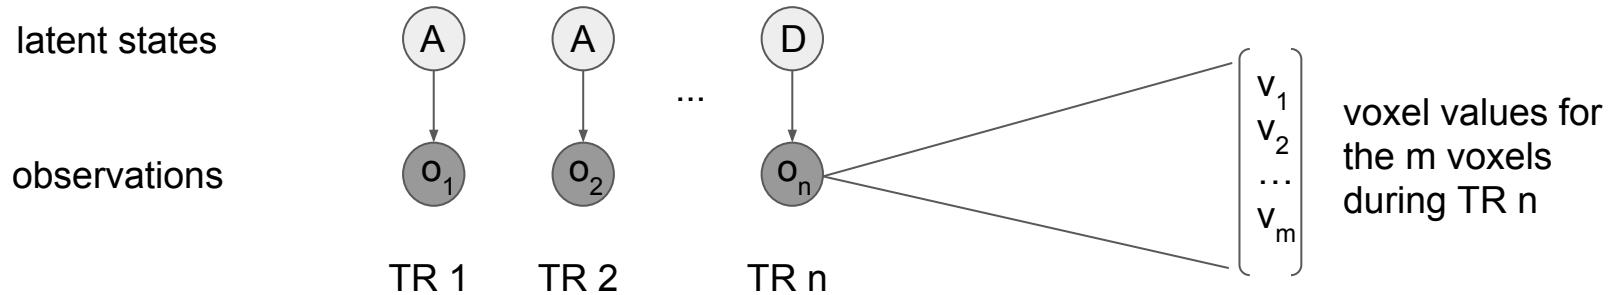
# What if there are latent processes?

- ❑ Dependence in observations may be due to dependence in latent states
  - ❑ Example:
    - ❑ underlying processes while subject is at rest in an fMRI
      - ❑ “attentive state” (A)
      - ❑ “default state” (D)



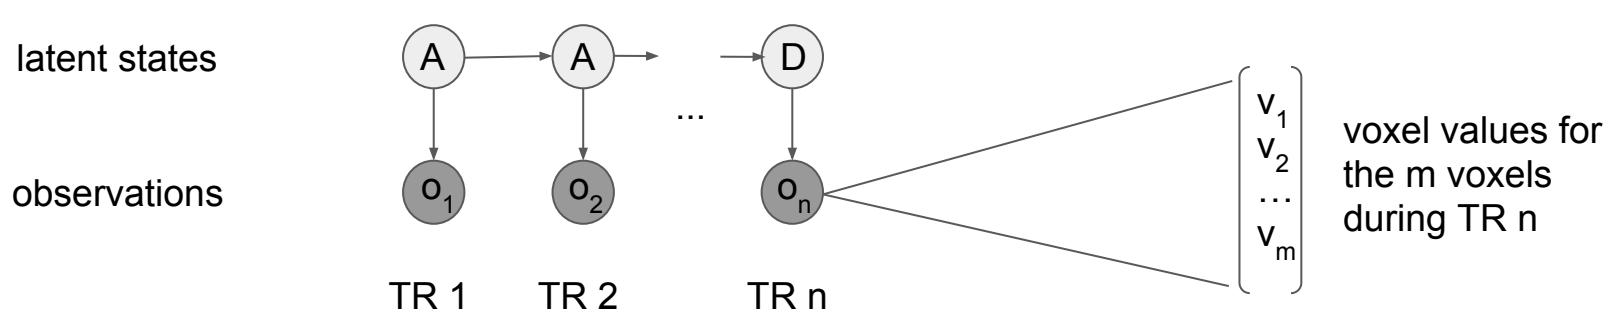
# What if there are latent processes?

- ❑ Dependence in observations may be due to dependence in latent states
  - ❑ Example:
    - ❑ underlying processes while subject is at rest in an fMRI
      - ❑ “attentive state” (A)
      - ❑ “default state” (D)
    - ❑ voxel values for ROIs in certain TRs may depend on whether subject is in state A or D

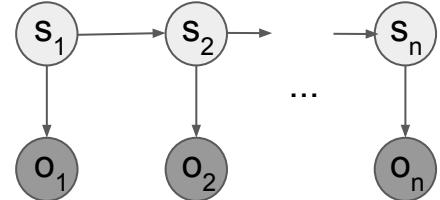


# What if there are latent processes?

- ❑ Dependence in observations may be due to dependence in latent states
  - ❑ Example:
    - ❑ underlying processes while subject is at rest in an fMRI
      - ❑ “attentive state” (A)
      - ❑ “default state” (D)
    - ❑ voxel values for ROIs in certain TRs may depend on whether subject is in state A or D
    - ❑ transitioning between states may help predict voxel values during next TR

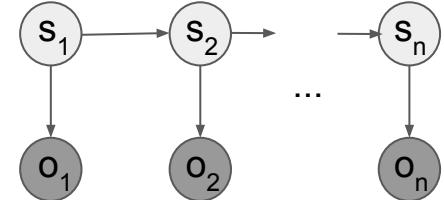


# Hidden Markov model (HMM)



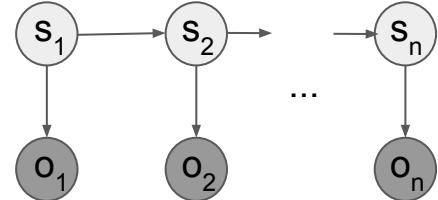
- ❑ Observation space  $O_t \in \{x_1, x_2, \dots, x_k\}$ 
  - ❑ observed data
  - ❑ can be discrete or real-valued
  - ❑ example: voxel values for ROI 1 during one TR

# Hidden Markov model (HMM)



- ❑ Observation space  $O_t \in \{x_1, x_2, \dots, x_k\}$ 
  - ❑ observed data
  - ❑ can be discrete or real-valued
  - ❑ example: voxel values for ROI 1 during one TR
- ❑ Hidden state space  $S_t \in \{1, \dots, m\}$ 
  - ❑ latent states
  - ❑ discrete
  - ❑ example: A (attentive state), D (default state)

# Hidden Markov model (HMM)



- ❑ Observation space  $O_t \in \{x_1, x_2, \dots, x_k\}$ 
  - ❑ observed data
  - ❑ can be discrete or real-valued
  - ❑ example: voxel values for ROI 1 during one TR
- ❑ Hidden state space  $S_t \in \{1, \dots, m\}$ 
  - ❑ latent states
  - ❑ discrete
  - ❑ example: A (attentive state), D (default state)
- ❑ Observed states are usually shaded

# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs



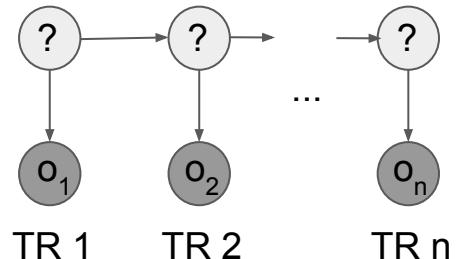
# What can we use HMMs for?

- We start with only having observations
  - e.g. voxel values for ROI 1 during all TRs
- Test a hypothesis about the underlying processes in the brain that lead to the observations



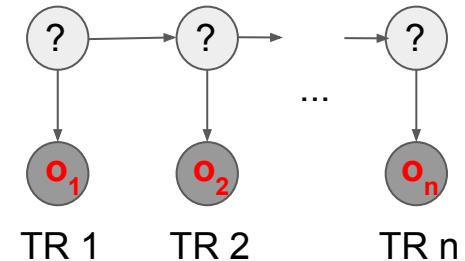
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state



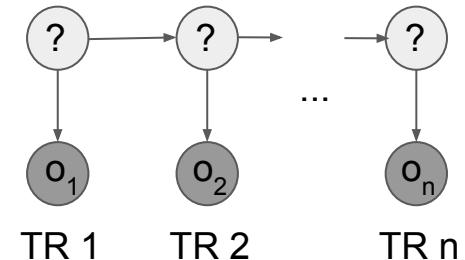
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state
  - ❑ test with HMM: given our hypothesis, how likely is the sequence of observed ROI voxel values  $\{o_1, \dots, o_n\}$ ?



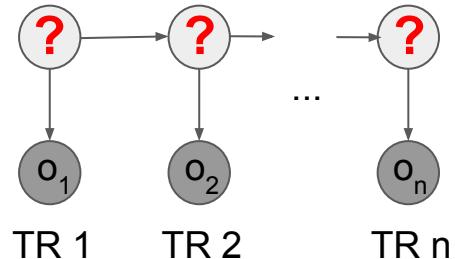
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state
  - ❑ test with HMM: given our hypothesis, how likely is the sequence of observed ROI voxel values  $\{o_1, \dots, o_n\}$ ?
- ❑ Infer the best sequence of hidden states that led to the observations



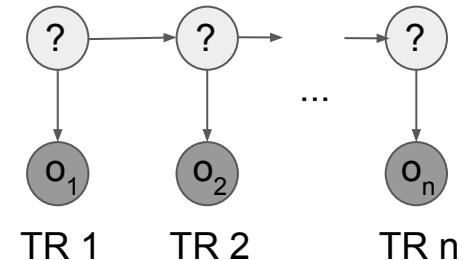
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state
  - ❑ test with HMM: given our hypothesis, how likely is the sequence of observed ROI voxel values  $\{o_1, \dots, o_n\}$ ?
- ❑ Infer the best sequence of hidden states that led to the observations
  - ❑ Inferring what brain state the subject was in during each TR



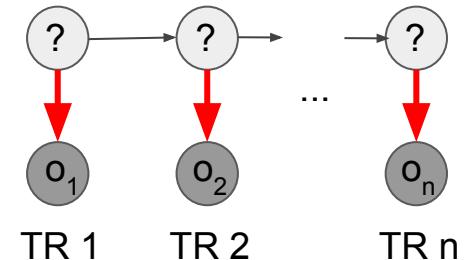
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state
  - ❑ test with HMM: given our hypothesis, how likely is the sequence of observed ROI voxel values  $\{o_1, \dots, o_n\}$ ?
- ❑ Infer the best sequence of hidden states that led to the observations
  - ❑ Inferring what brain state the subject was in during each TR
- ❑ How are observations emitted from the hidden states? How does the model transition from one hidden state to another?



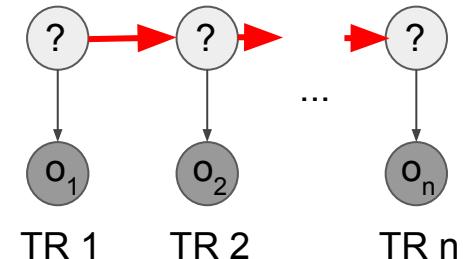
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state
  - ❑ test with HMM: given our hypothesis, how likely is the sequence of observed ROI voxel values  $\{o_1, \dots, o_n\}$ ?
- ❑ Infer the best sequence of hidden states that led to the observations
  - ❑ Inferring what brain state the subject was in during each TR
- ❑ How are observations emitted from the hidden states? How does the model transition from one hidden state to another?
  - ❑ Learn how the observed voxel values for a TR are related to the current brain state



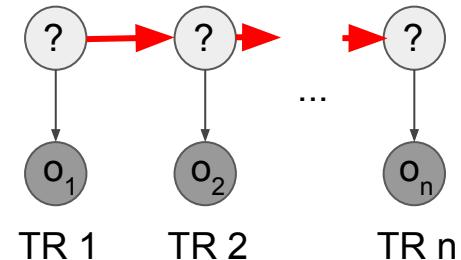
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state
  - ❑ test with HMM: given our hypothesis, how likely is the sequence of observed ROI voxel values  $\{o_1, \dots, o_n\}$ ?
- ❑ Infer the best sequence of hidden states that led to the observations
  - ❑ Inferring what brain state the subject was in during each TR
- ❑ How are observations emitted from the hidden states? How does the model transition from one hidden state to another?
  - ❑ Learn how the observed voxel values for a TR are related to the current brain state
  - ❑ Learn how the brain transitions between attentive and default states



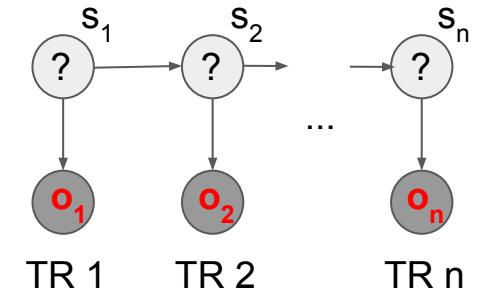
# What can we use HMMs for?

- ❑ We start with only having observations
  - ❑ e.g. voxel values for ROI 1 during all TRs
- ❑ Test a hypothesis about the underlying processes in the brain that lead to the observations   **Forward algorithm**
  - ❑ hypothesis: at rest, the brain switches between an attentive and default state
  - ❑ test with HMM: given our hypothesis, how likely is the sequence of observed ROI voxel values  $\{o_1, \dots, o_n\}$ ?
- ❑ Infer the best sequence of hidden states that led to the observations
  - ❑ Inferring what brain state the subject was in during each TR   **Viterbi algorithm**
- ❑ How are observations emitted from the hidden states? How does the model transition from one hidden state to another?   **Expectation-Maximization (EM)**
  - ❑ Learn how the observed voxel values for a TR are related to the current brain state
  - ❑ Learn how the brain transitions between attentive and default states



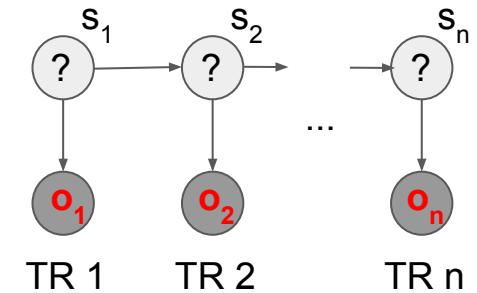
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)



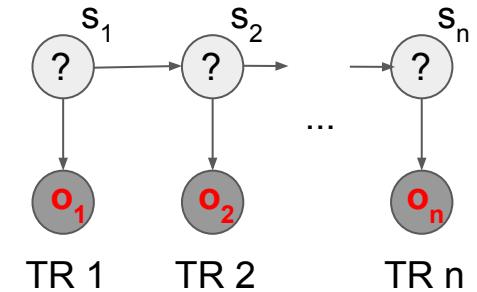
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM



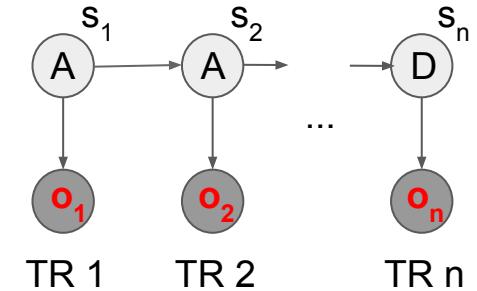
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model



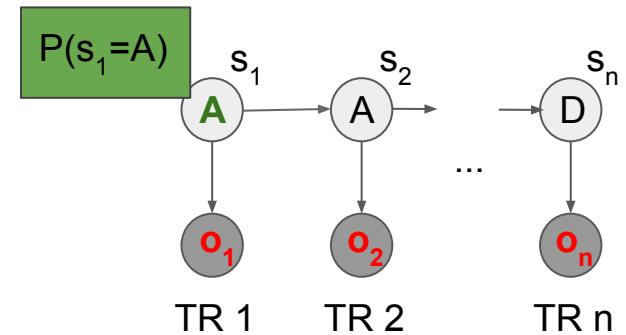
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$



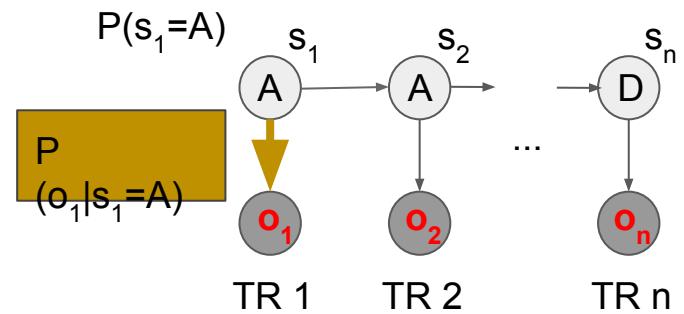
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$



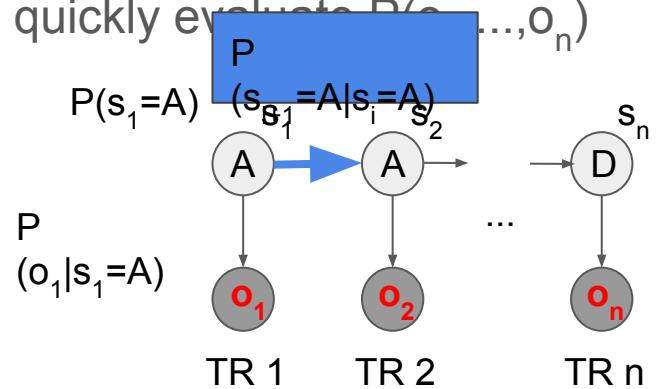
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$



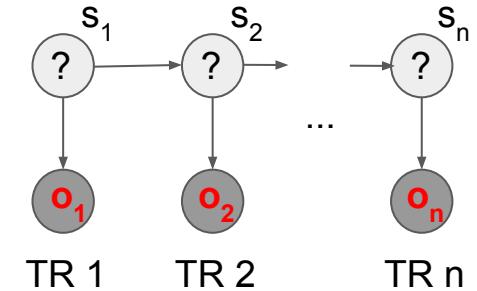
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$



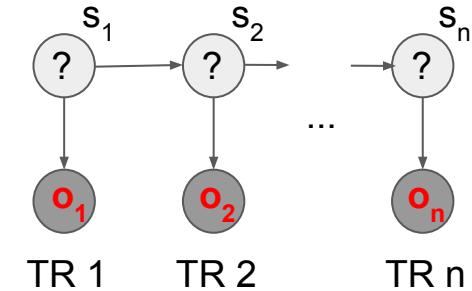
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$
- ❑ But we don't know what the hidden state values are!



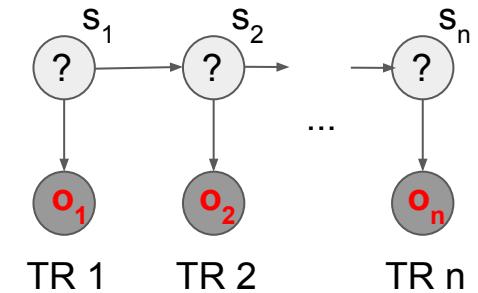
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$
- ❑ But we don't know what the hidden state values are!
- ❑ So how can we evaluate  $P(o_1, \dots, o_n)$ ?



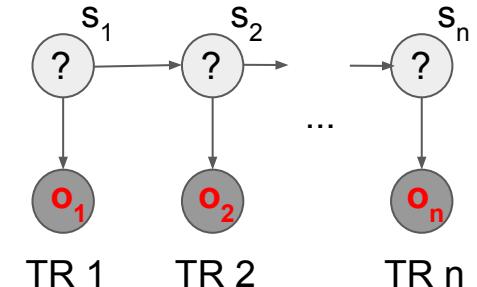
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$
- ❑ But we don't know what the hidden state values are!
- ❑ So how can we evaluate  $P(o_1, \dots, o_n)$ ?
  - ❑ calculate  $P(o_1, \dots, o_n)$  under all possibilities for  $\{s_1, \dots, s_n\}$  and sum them since we don't know which  $\{s_1, \dots, s_n\}$  is correct



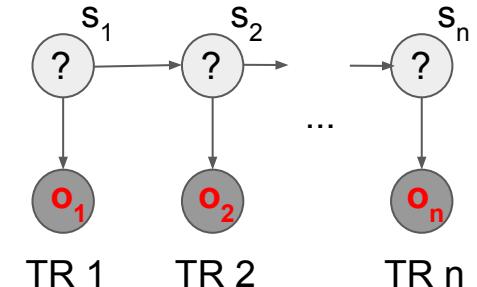
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$
- ❑ But we don't know what the hidden state values are!
- ❑ So how can we evaluate  $P(o_1, \dots, o_n)$ ?
  - ❑ calculate  $P(o_1, \dots, o_n)$  under all possibilities for  $\{s_1, \dots, s_n\}$  and sum them since we don't know which  $\{s_1, \dots, s_n\}$  is correct
- ❑ How many possibilities for  $\{s_1, \dots, s_n\}$ ?



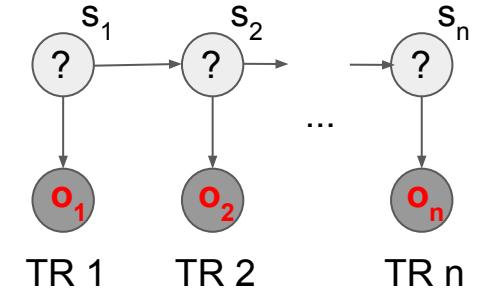
# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$
- ❑ But we don't know what the hidden state values are!
- ❑ So how can we evaluate  $P(o_1, \dots, o_n)$ ?
  - ❑ calculate  $P(o_1, \dots, o_n)$  under all possibilities for  $\{s_1, \dots, s_n\}$  and sum them since we don't know which  $\{s_1, \dots, s_n\}$  is correct
- ❑ How many possibilities for  $\{s_1, \dots, s_n\}$ ?  $2^n$



# How likely is the sequence of observed ROI voxel values $\{o_1, \dots, o_n\}$ ?

- ❑ Hidden states  $\{s_1, \dots, s_n\}$  can take 2 values: A (attentive) or D (default)
- ❑ Assume that we already learned both the emissions  $P(o_i|s_i)$  and transitions  $P(s_{i+1}|s_i)$  with EM
- ❑ Want to evaluate  $P(o_1, \dots, o_n)$  under this model
- ❑ If we know what the values of  $s_1, \dots, s_n$  are, we quickly evaluate  $P(o_1, \dots, o_n)$ 
  - ❑ example: if  $s_1 = A, s_2 = A, \dots, s_3 = D$
- ❑ But we don't know what the hidden state values are!
- ❑ So how can we evaluate  $P(o_1, \dots, o_n)$ ?
  - ❑ calculate  $P(o_1, \dots, o_n)$  under all possibilities for  $\{s_1, \dots, s_n\}$  and sum them since we don't know which  $\{s_1, \dots, s_n\}$  is correct
- ❑ How many possibilities for  $\{s_1, \dots, s_n\}$ ?  $2^n$ 
  - ❑ forward algorithm is a clever way to do this

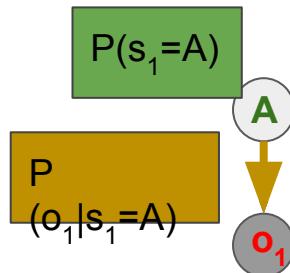


# Forward algorithm: computationally friendly way to calculate probability of observed sequence

- ❑ Key idea: recognize that there is a lot of overlap between possibilities for  $\{s_1, \dots, s_n\}$ , so instead of recalculating the whole probability for each possibility, store probabilities for smaller pieces of the sequence and reuse

# Forward algorithm: computationally friendly way to calculate probability of observed sequence

- ❑ Key idea: recognize that there is a lot of overlap between possibilities for  $\{s_1, \dots, s_n\}$ , so instead of recalculating the whole probability for each possibility, store probabilities for smaller pieces of the sequence and reuse
- ❑ First subunits:  $P(s_1=A)P(o_1|s_1=A)$  and  $P(s_1=D)P(o_1|s_1=D)$



# Forward algorithm: computationally friendly way to calculate probability of observed sequence

- ❑ Key idea: recognize that there is a lot of overlap between possibilities for  $\{s_1, \dots, s_n\}$ , so instead of recalculating the whole probability for each possibility, store probabilities for smaller pieces of the sequence and reuse
- ❑ First subunits:  $P(s_1=A)P(o_1|s_1=A)$  and  $P(s_1=D)P(o_1|s_1=D)$
- ❑ All further subunits: sum of **emission** x **transition** x previous subunit



## Problem 2: infer best sequence of hidden states that results in the sequence of observations

- ❑ Viterbi algorithm: also relies on recursive computation using subunits

## Problem 2: infer best sequence of hidden states that results in the sequence of observations

- ❑ Viterbi algorithm: also relies on recursive computation using subunits
- ❑ Starting with the second hidden state, for each hidden state until the last one, we calculate the probability of the most likely path that ends in this hidden state

# Problem 2: infer best sequence of hidden states that results in the sequence of observations

- ❑ Viterbi algorithm: also relies on recursive computation using subunits
- ❑ Starting with the second hidden state, for each hidden state until the last one, we calculate the probability of the most likely path that ends in this hidden state
- ❑ For each of these probabilities, we record the previous hidden state in the most likely sequence => this enables us to follow the most probable path backwards once we're at the end of the sequence

# Problem 3: learning HMM emissions and transitions

- ❑ So far, we've assumed that the emission and transition probabilities were already learned

# Problem 3: learning HMM emissions and transitions

- ❑ So far, we've assumed that the emission and transition probabilities were already learned
- ❑ But how do we do this? Expectation-maximization (EM)!

# Problem 3: learning HMM emissions and transitions

- ❑ So far, we've assumed that the emission and transition probabilities were already learned
- ❑ But how do we do this? Expectation-maximization (EM)!
- ❑ Start with random initializations of parameters, then iterate 2 steps:

# Problem 3: learning HMM emissions and transitions

- ❑ So far, we've assumed that the emission and transition probabilities were already learned
- ❑ But how do we do this? Expectation-maximization (EM)!
- ❑ Start with random initializations of parameters, then iterate 2 steps:
  - ❑ E-step: fix parameters, find expected state assignments

# Problem 3: learning HMM emissions and transitions

- ❑ So far, we've assumed that the emission and transition probabilities were already learned
- ❑ But how do we do this? Expectation-maximization (EM)!
- ❑ Start with random initializations of parameters, then iterate 2 steps:
  - ❑ E-step: fix parameters, find expected state assignments
  - ❑ M-step: fix expected state assignments, update parameters

# Problem 3: learning HMM emissions and transitions

- ❑ So far, we've assumed that the emission and transition probabilities were already learned
- ❑ But how do we do this? Expectation-maximization (EM)!
- ❑ Start with random initializations of parameters, then iterate 2 steps:
  - ❑ E-step: fix parameters, find expected state assignments
  - ❑ M-step: fix expected state assignments, update parameters
- ❑ Very general algorithm for learning parameters in models with latent variables

# Latent variable models takeaways

- ❑ Used to analyze sequential data (such as time series)

# Latent variable models takeaways

- ❑ Used to analyze sequential data (such as time series)
- ❑ Hidden Markov model is an example of such a model, where the hidden states can take discrete values

# Latent variable models takeaways

- ❑ Used to analyze sequential data (such as time series)
- ❑ Hidden Markov model is an example of such a model, where the hidden states can take discrete values
- ❑ We can use HMMs for 3 main purposes:

# Latent variable models takeaways

- ❑ Used to analyze sequential data (such as time series)
- ❑ Hidden Markov model is an example of such a model, where the hidden states can take discrete values
- ❑ We can use HMMs for 3 main purposes:
  - ❑ Testing a hypothesis about underlying processes in the brain

# Latent variable models takeaways

- ❑ Used to analyze sequential data (such as time series)
- ❑ Hidden Markov model is an example of such a model, where the hidden states can take discrete values
- ❑ We can use HMMs for 3 main purposes:
  - ❑ Testing a hypothesis about underlying processes in the brain
  - ❑ Inferring the best sequence of underlying processes in the brain that emits the observations

# Latent variable models takeaways

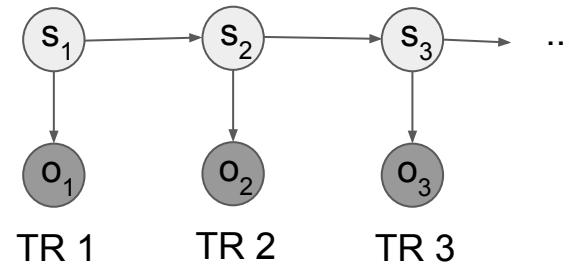
- ❑ Used to analyze sequential data (such as time series)
- ❑ Hidden Markov model is an example of such a model, where the hidden states can take discrete values
- ❑ We can use HMMs for 3 main purposes:
  - ❑ Testing a hypothesis about underlying processes in the brain
  - ❑ Inferring the best sequence of underlying processes in the brain that emits the observations
  - ❑ Learning how the observations were emitted from the brain states, and how the brain transitions from one underlying state to another

# Latent variable models takeaways

- ❑ Used to analyze sequential data (such as time series)
- ❑ Hidden Markov model is an example of such a model, where the hidden states can take discrete values
- ❑ We can use HMMs for 3 main purposes:
  - ❑ Testing a hypothesis about underlying processes in the brain
  - ❑ Inferring the best sequence of underlying processes in the brain that emits the observations
  - ❑ Learning how the observations were emitted from the brain states, and how the brain transitions from one underlying state to another
- ❑ SSM (state-space model) is another type of latent variable model that accounts for continuous hidden states

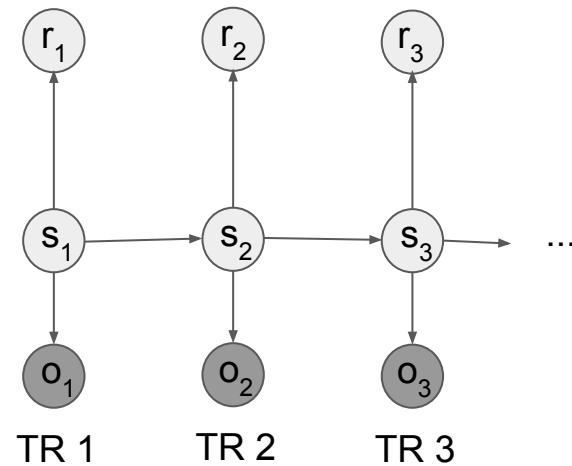
# Can we model human learning with an HMM?

- ❑ To learn, we need some teaching, or reward, signal



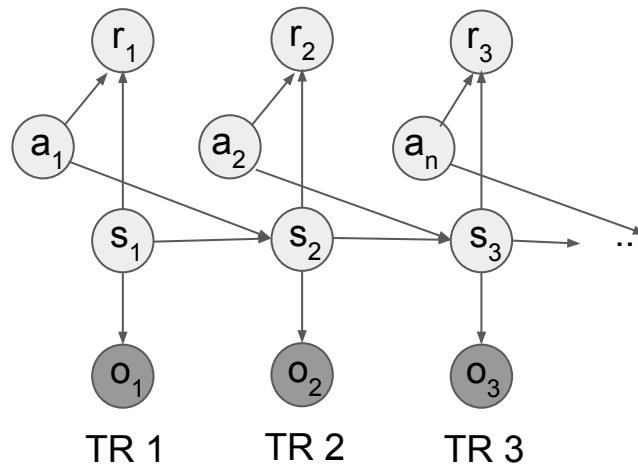
# Can we model human learning with an HMM?

- ❑ To learn, we need some teaching, or reward, signal
- ❑ What if we incorporate a notion of reward in the HMM?



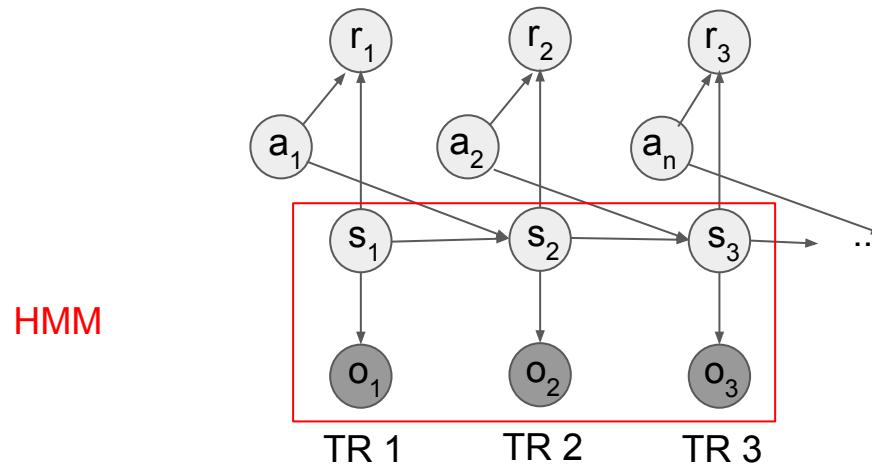
# Can we model human learning with an HMM?

- ❑ To learn, we need some teaching, or reward, signal
- ❑ What if we incorporate a notion of reward in the HMM?
- ❑ Now, we incorporate actions we can take to transition to other states in order to maximize our reward



# Can we model human learning with an HMM?

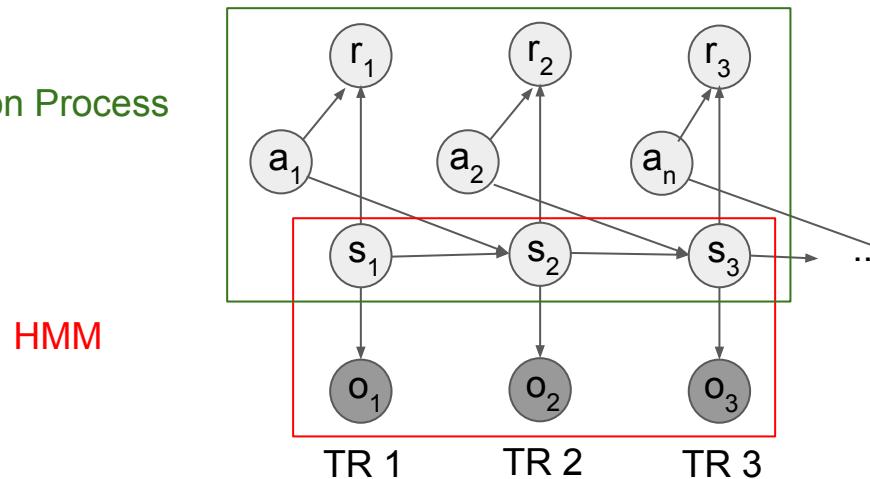
- ❑ To learn, we need some teaching, or reward, signal
- ❑ What if we incorporate a notion of reward in the HMM?
- ❑ Now, we incorporate actions we can take to transition to other states in order to maximize our reward



# Can we model human learning with an HMM?

- ❑ To learn, we need some teaching, or reward, signal
- ❑ What if we incorporate a notion of reward in the HMM?
- ❑ Now, we incorporate actions we can take to transition to other states in order to maximize our reward

Markov Decision Process  
(MDP)



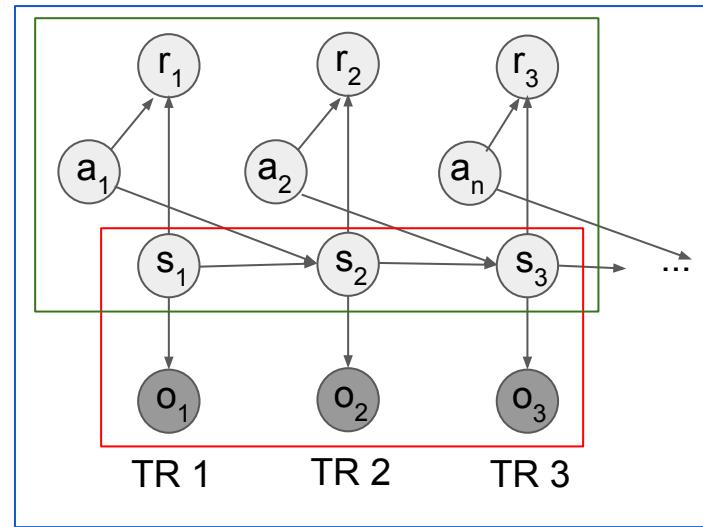
# Can we model human learning with an HMM?

- ❑ To learn, we need some teaching, or reward, signal
- ❑ What if we incorporate a notion of reward in the HMM?
- ❑ Now, we incorporate actions we can take to transition to other states in order to maximize our reward

Markov Decision Process  
(MDP)

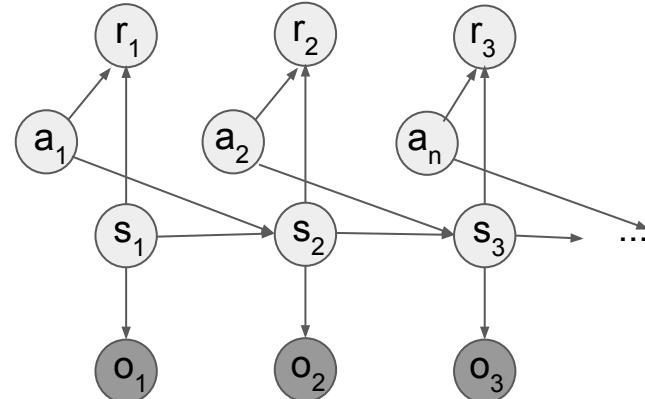
HMM

Partially Observable  
Markov Decision Process  
(POMDP)



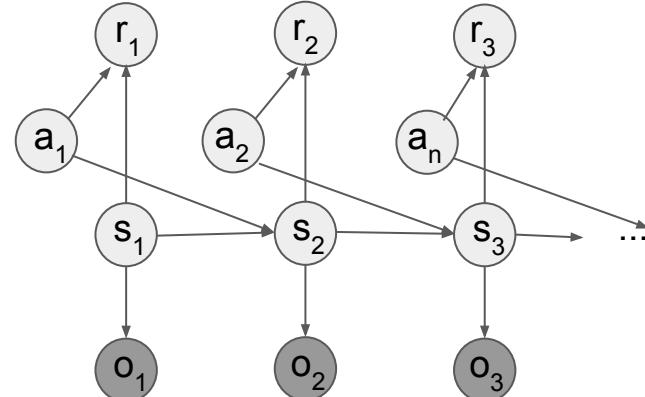
# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward



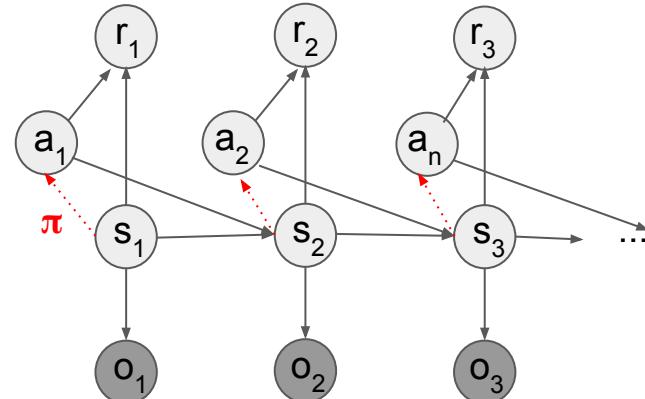
# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods



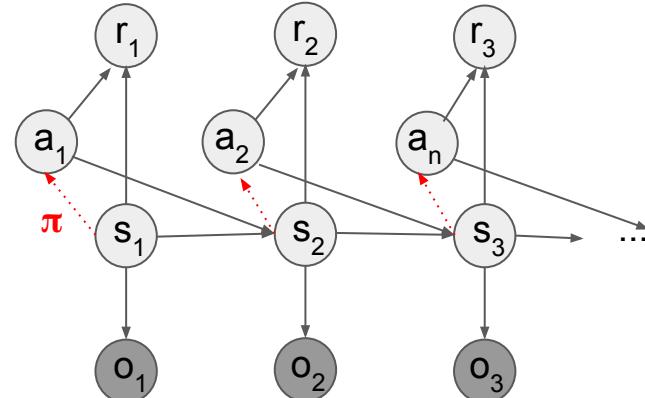
# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi: S \rightarrow A$  that tells us what the best action to take is in the current state



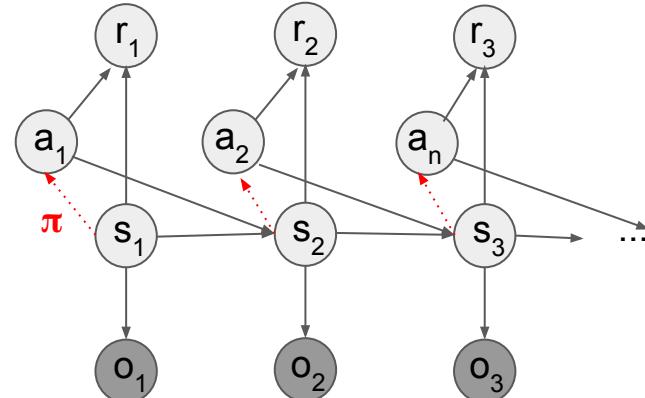
# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi: S \rightarrow A$  that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”



# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

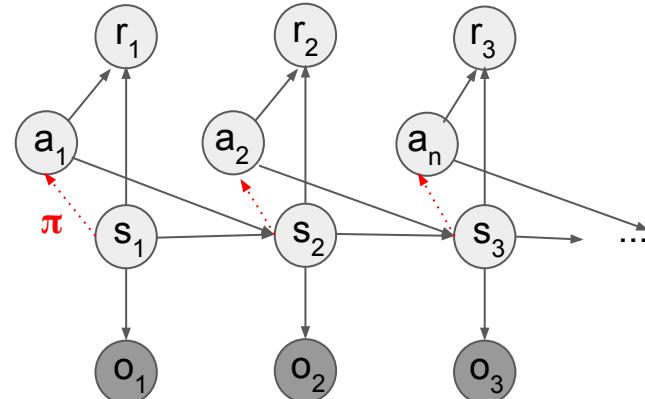
- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi: S \rightarrow A$  that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”
  - ❑ Given any policy  $\pi$ , we can define the expected future reward if we **start at state  $s$**  as:



# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi$ : S->A that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”
  - ❑ Given any policy  $\pi$ , we can define the expected future reward if we **start at state s** as:

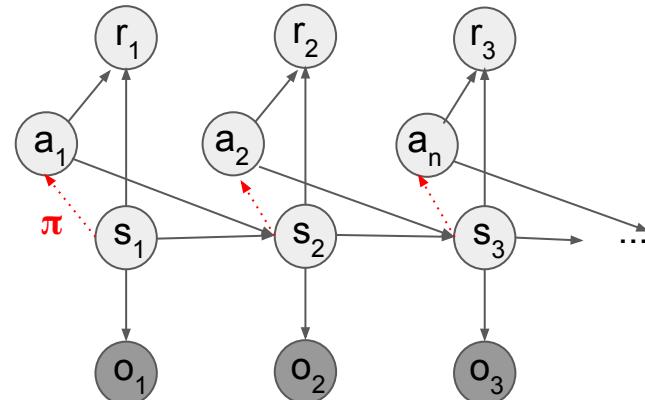
$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$



# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi$ : S->A that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”
  - ❑ Given any policy  $\pi$ , we can define the expected future **reward** if we start at state s as:

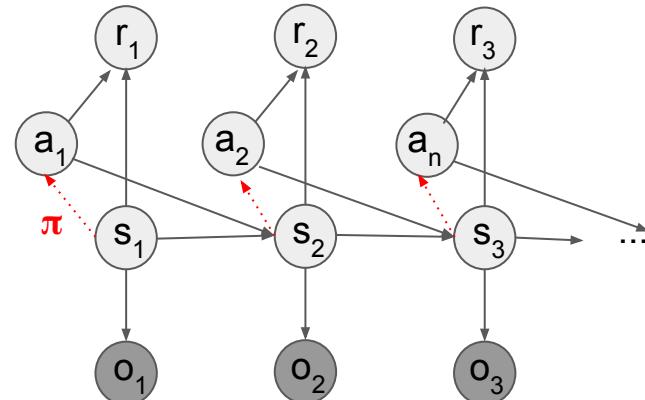
$$V^\pi(s) = E\left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$



# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

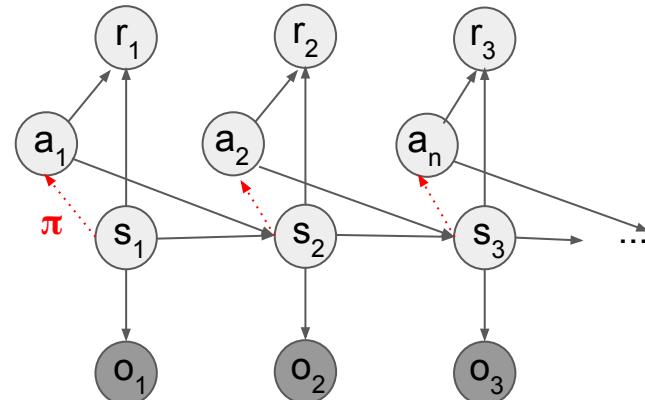
- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi$ : S->A that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”
  - ❑ Given any policy  $\pi$ , we can define the expected **future** reward if we start at state s as:

$$V^\pi(s) = E\left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$



# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi$ : S->A that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”
  - ❑ Given any policy  $\pi$ , we can define the expected **future** reward if we start at state s as:
$$V^\pi(s) = E\left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$
  - ❑ We call this the value function of policy  $\pi$

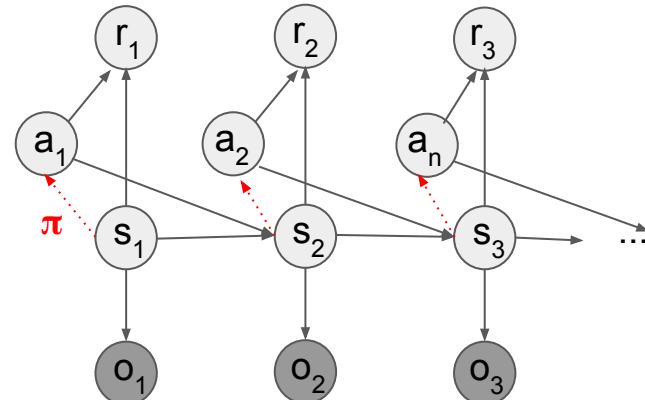


# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi$ : S->A that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”
  - ❑ Given any policy  $\pi$ , we can define the expected **future** reward if we start at state s as:

$$V^\pi(s) = E\left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

- ❑ We call this the value function of policy  $\pi$
- ❑ Now, we can compute the best policy:



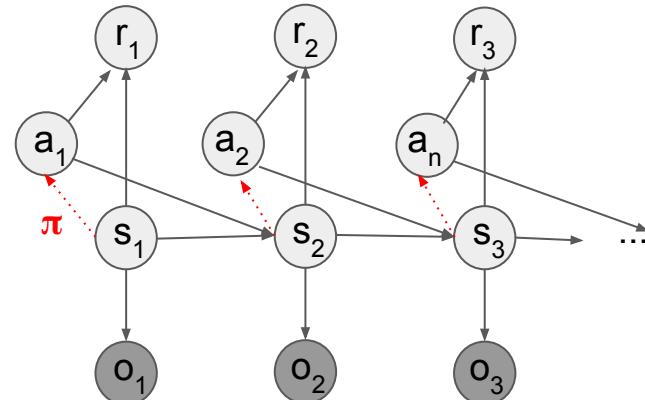
# Reinforcement learning (RL): learning how to take the best actions to maximize expected future reward

- ❑ Delayed reward vs immediate reward
  - ❑ We can't choose the best action to take right now with regular supervised methods
- ❑ RL in a nutshell:
  - ❑ Learning function  $\pi: S \rightarrow A$  that tells us what the best action to take is in the current state
    - ❑  $\pi$  is also called a “policy”
  - ❑ Given any policy  $\pi$ , we can define the expected **future** reward if we start at state  $s$  as:

$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

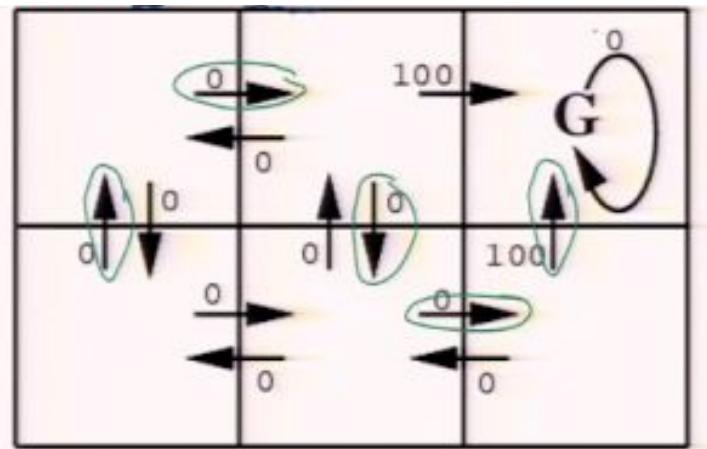
- ❑ We call this the value function of policy  $\pi$
- ❑ Now, we can compute the best policy:

$$\pi^* = \arg \max_{\pi} V^\pi(s), \quad (\forall s)$$



# Example: what are the $V^\pi(s)$ values?

- ❑ Image there is a robot that needs to get to the goal state G
- ❑ Each square is a state
- ❑ Circled arrows show the policy  $\pi$
- ❑ Let  $\gamma = 0.9$

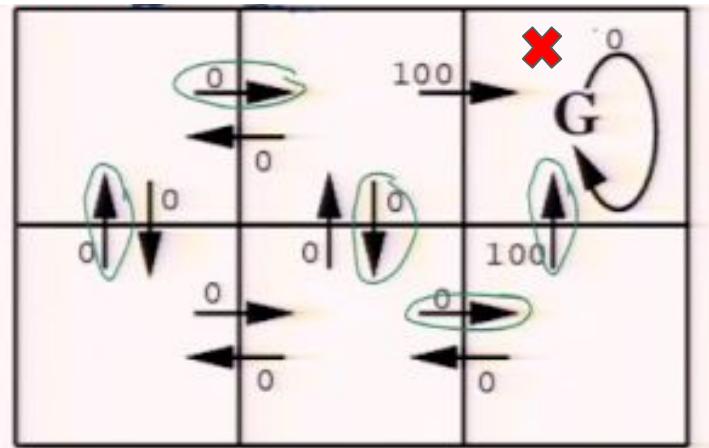


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

What is  $V^\pi(\text{x})$ ?

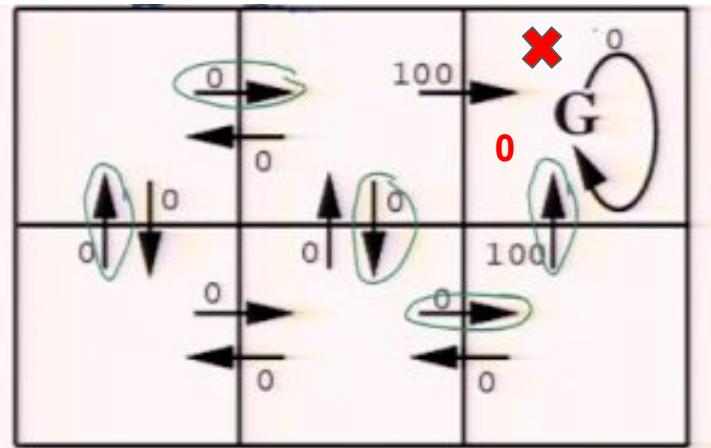


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

What is  $V^\pi(\text{red } X)$ ? 0

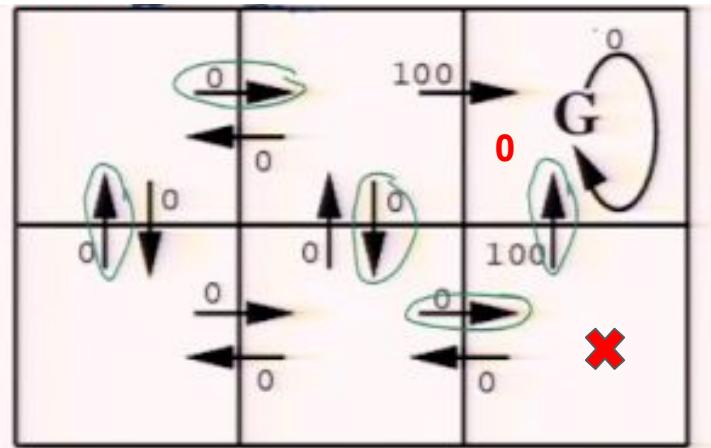


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

What is  $V^\pi(\text{x})$ ?

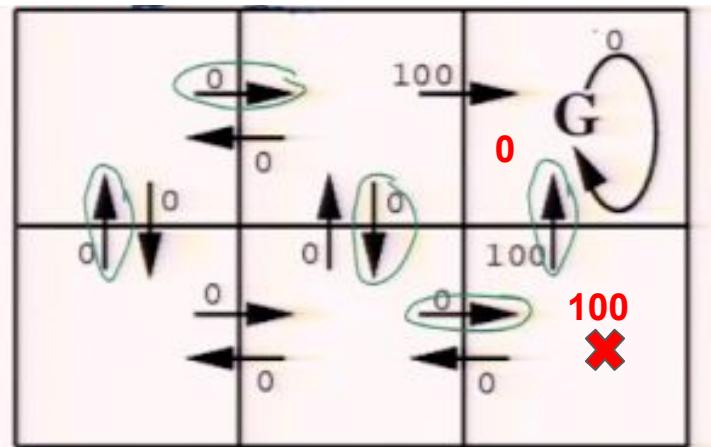


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

What is  $V^\pi(x)$ ? 100

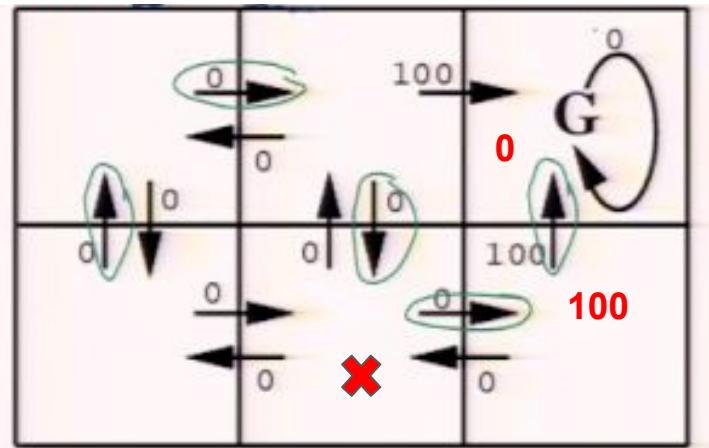


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

What is  $V^\pi(\text{x})$ ?

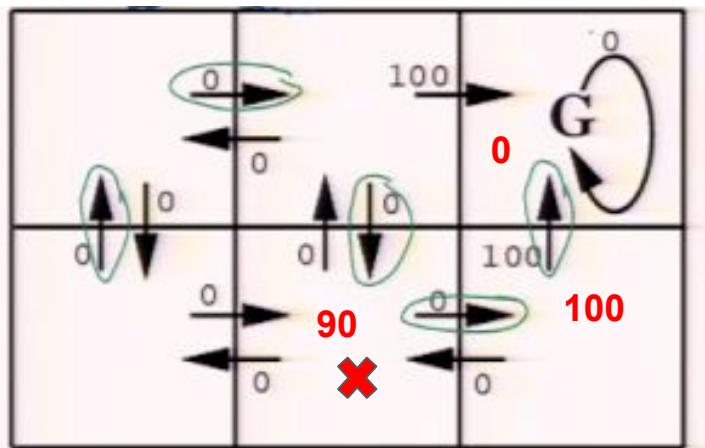


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

What is  $V^\pi(x)$ ? 0.9  
 $\times 100 = 90$

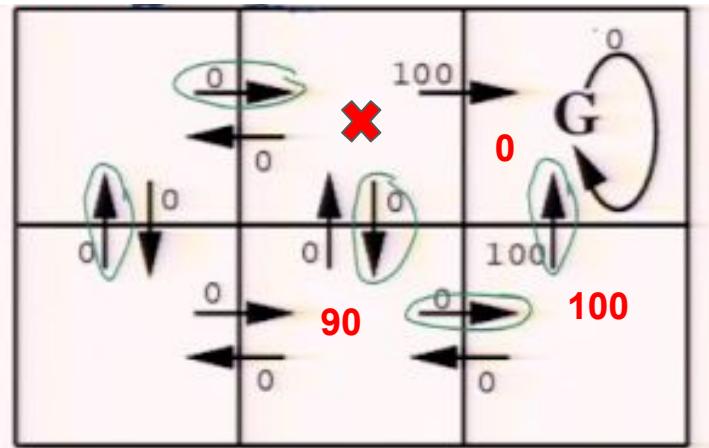


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

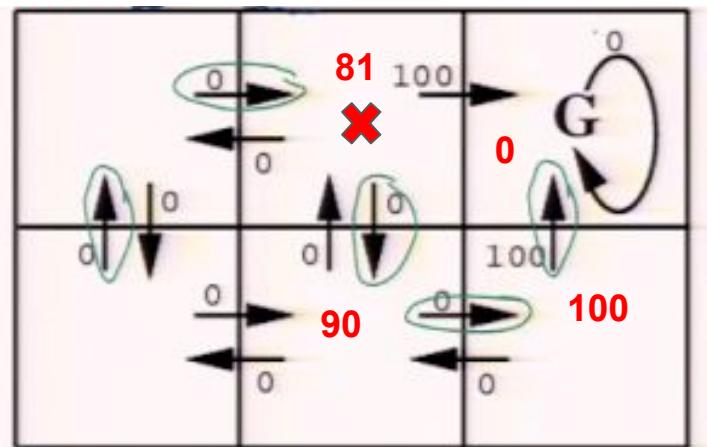
What is  $V^\pi(x)$ ?



$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$



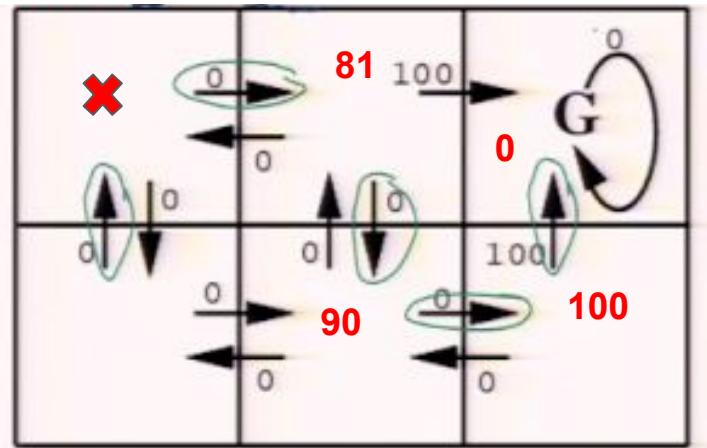
What is  $V^\pi(x)$ ?  $0.9 \times 0$ .  
 $9 \times 100 = 81$

$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

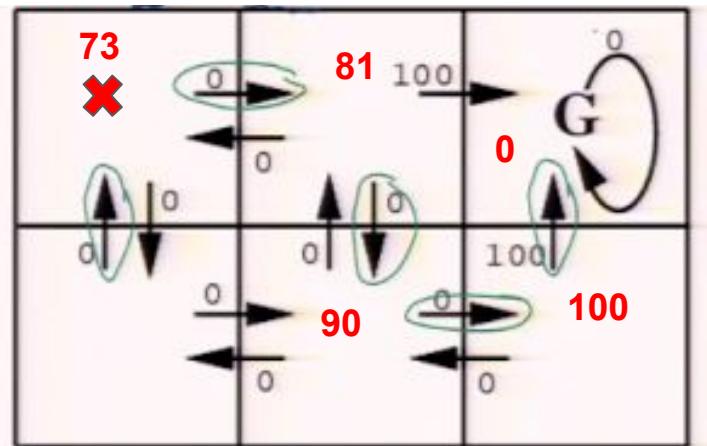
What is  $V^\pi(x)$ ?



$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$



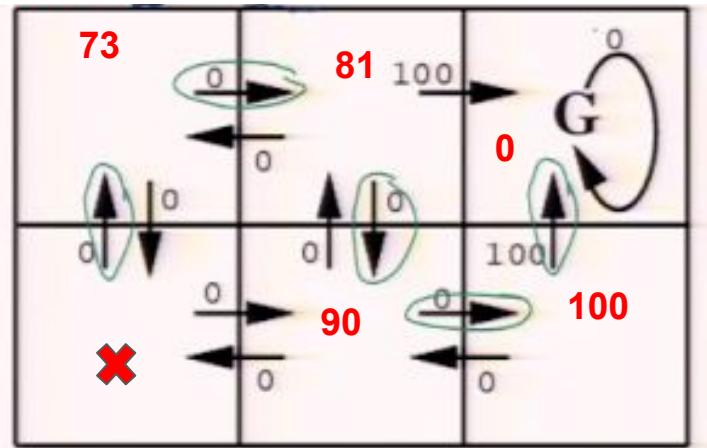
What is  $V^\pi(\text{X})$ ?  $0.9 \times 0$ .  
 $9 \times 0.9 \times 100 = 73$

$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

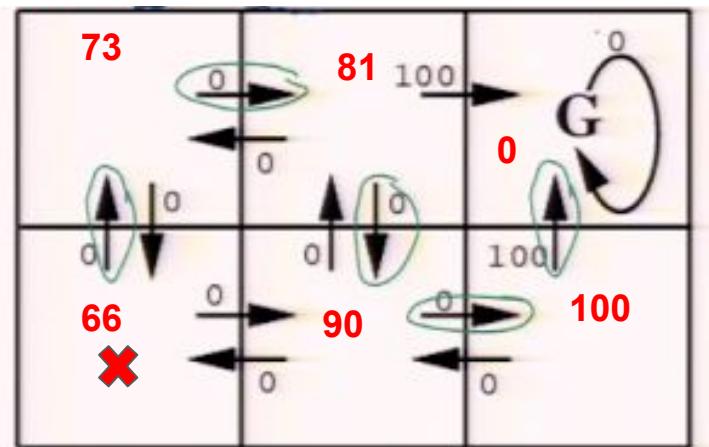
What is  $V^\pi(x)$ ?



$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the $V^\pi(s)$ values?

- Image there is a robot that needs to get to the goal state G
- Each square is a state
- Circled arrows show the policy  $\pi$
- Let  $\gamma = 0.9$

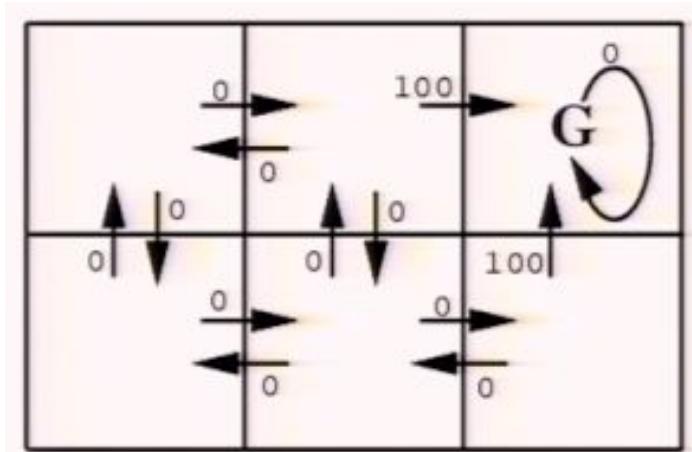


What is  $V^\pi(x)$ ?  $0.9 \times 0.$   
 $9 \times 0.9 \times 100 = 66$

$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

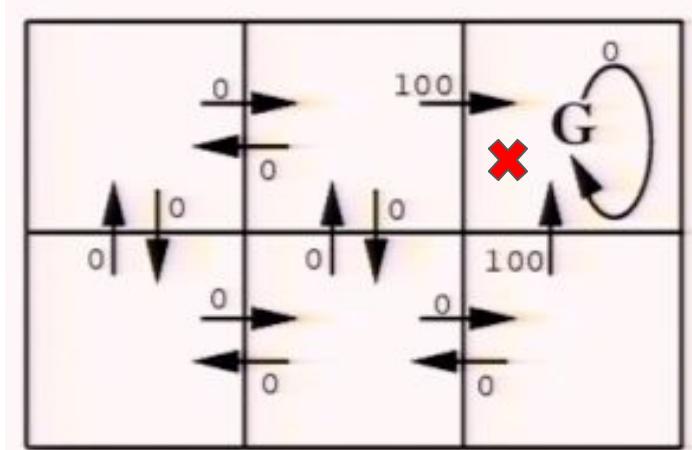


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\text{x})$ ?

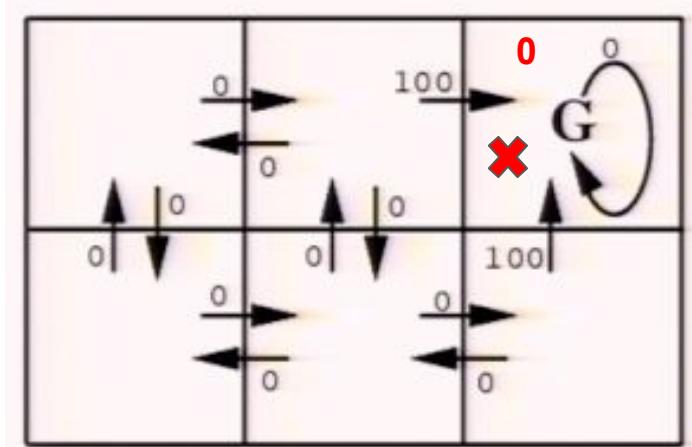


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\text{x})$ ? Still 0

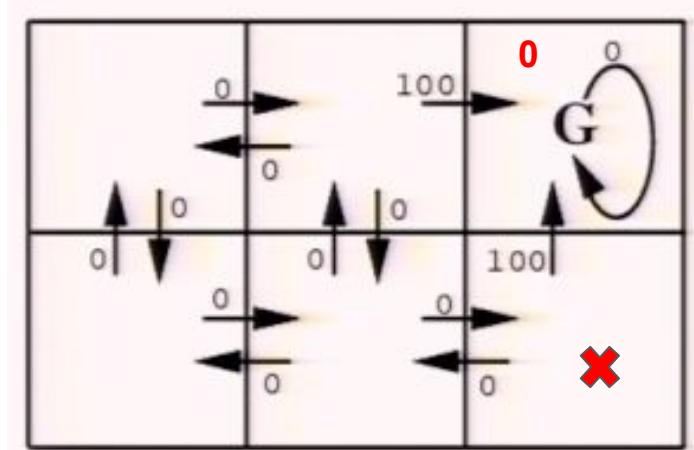


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\text{x})$ ?

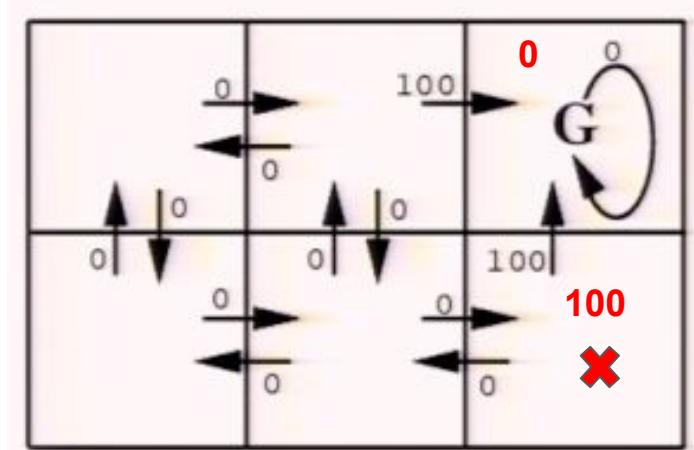


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\text{x})$ ? 100

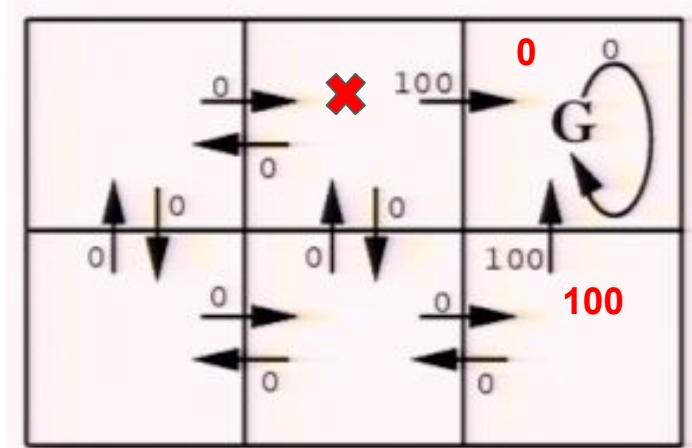


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\mathbf{x})$ ?

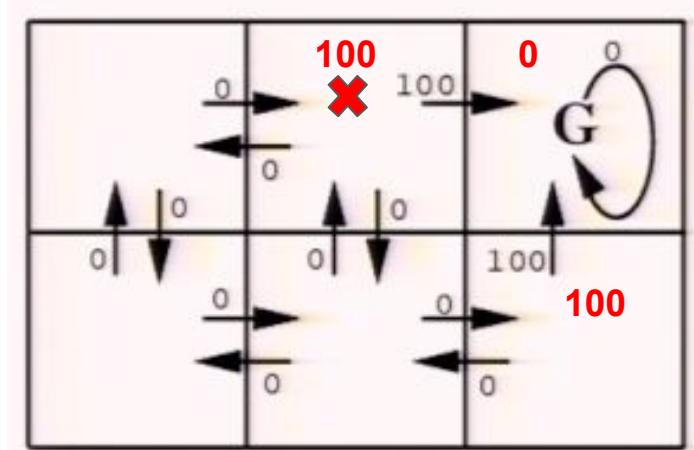


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\mathbf{x})$ ? 100

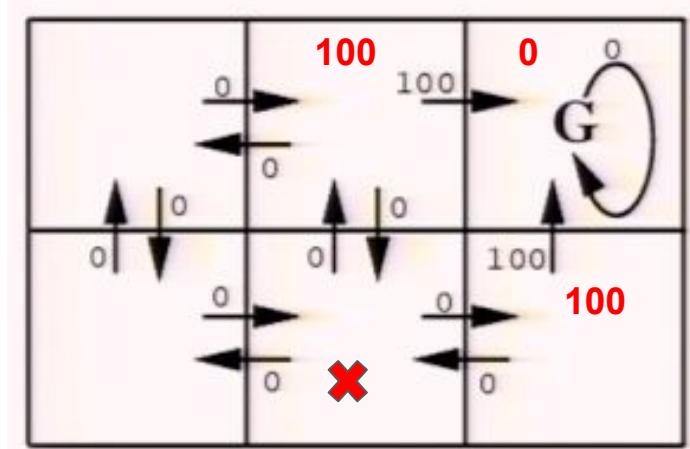


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\text{x})$ ?

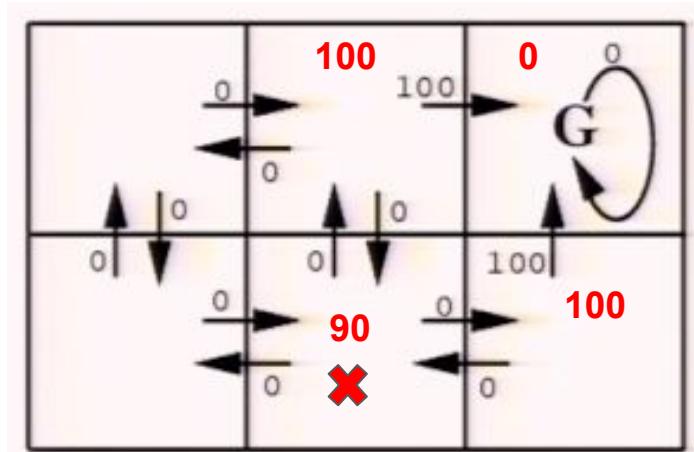


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\text{x})$ ? 90

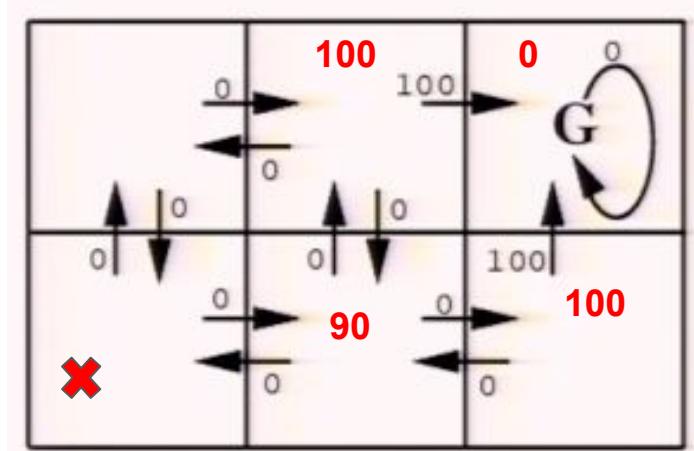


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\text{x})$ ?

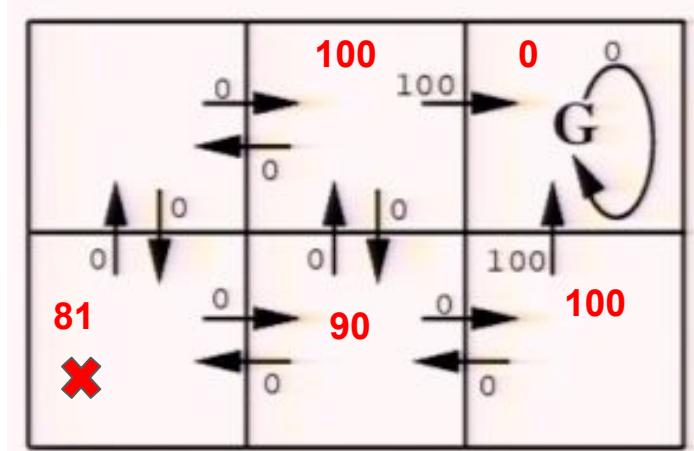


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(x)$ ? 81

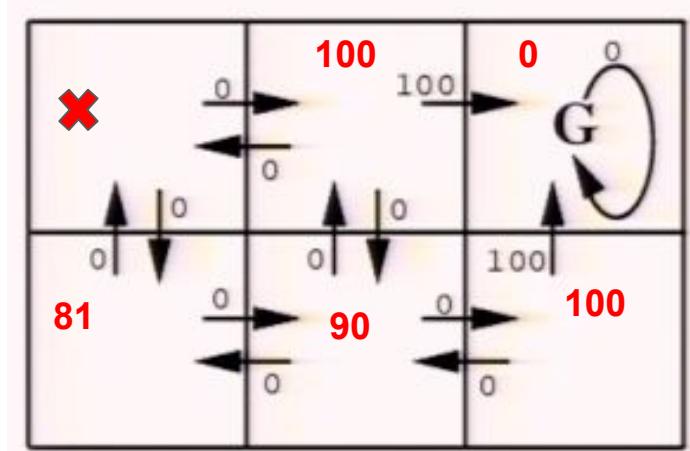


$$V^\pi(s) = E\left[\sum_{t=0}^{\infty} \gamma^t r_t\right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\mathbf{x})$ ?

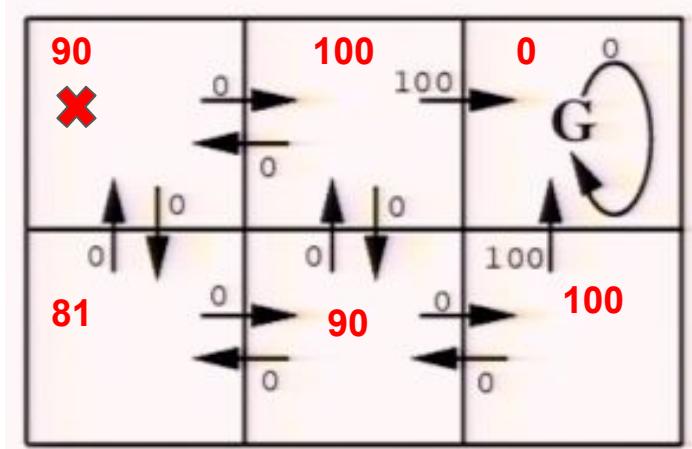


$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# Example: what are the best $V^*(s)$ values?

- What is the best policy?

What is  $V^*(\mathbf{x})$ ? 90



$$V^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t r_t \right]$$

# RL takeaways

- ❑ Reinforcement learning introduces the notion of delayed reward and uses it to learn the best actions that will maximize this future reward

# RL takeaways

- ❑ Reinforcement learning introduces the notion of delayed reward and uses it to learn the best actions that will maximize this future reward
- ❑ Two ways to learn the best policy:

# RL takeaways

- ❑ Reinforcement learning introduces the notion of delayed reward and uses it to learn the best actions that will maximize this future reward
- ❑ Two ways to learn the best policy:
  - ❑ Q-learning
  - ❑ Temporal Difference (TD) learning

# RL takeaways

- ❑ Reinforcement learning introduces the notion of delayed reward and uses it to learn the best actions that will maximize this future reward
- ❑ Two ways to learn the best policy:
  - ❑ Q-learning
  - ❑ Temporal Difference (TD) learning
    - ❑ Error in value prediction using TD-learning is the best computational model of dopamine release in monkeys during anticipation of future reward => hypothesis that dopamine does reward prediction (Schultz et al. Science, 1997)

# RL takeaways

- ❑ Reinforcement learning introduces the notion of delayed reward and uses it to learn the best actions that will maximize this future reward
- ❑ Two ways to learn the best policy:
  - ❑ Q-learning
  - ❑ Temporal Difference (TD) learning
    - ❑ Error in value prediction using TD-learning is the best computational model of dopamine release in monkeys during anticipation of future reward => hypothesis that dopamine does reward prediction (Schultz et al. Science, 1997)
- ❑ Recent advances in RL: combining it with deep learning!

# RL takeaways

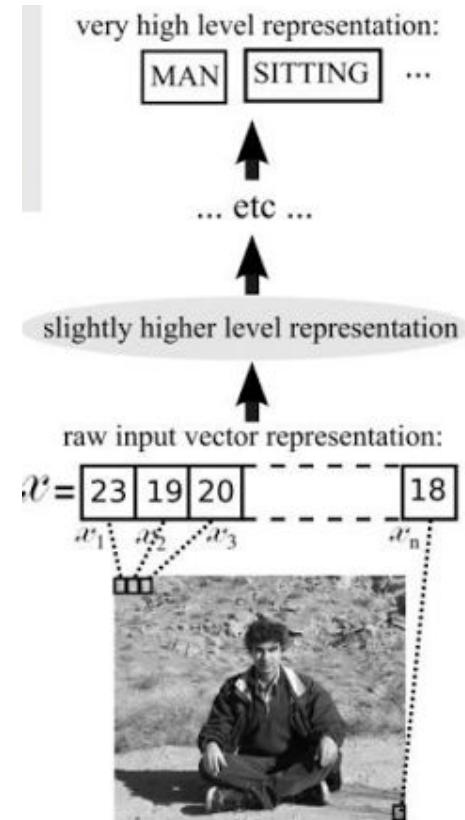
- ❑ Reinforcement learning introduces the notion of delayed reward and uses it to learn the best actions that will maximize this future reward
- ❑ Two ways to learn the best policy:
  - ❑ Q-learning
  - ❑ Temporal Difference (TD) learning
    - ❑ Error in value prediction using TD-learning is the best computational model of dopamine release in monkeys during anticipation of future reward => hypothesis that dopamine does reward prediction (Schultz et al. Science, 1997)
- ❑ Recent advances in RL: combining it with deep learning!
  - ❑ learning deep policy and value networks

# Deep learning!

- ❑ Loosely inspired by neural connections in the brain

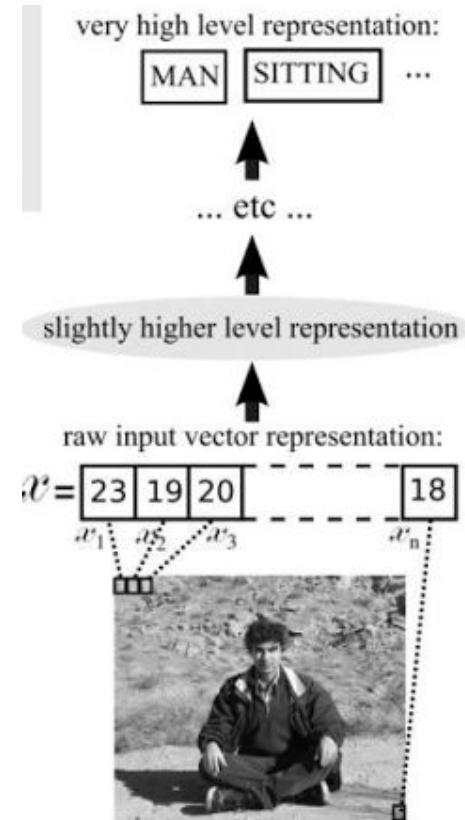
# Deep learning!

- ❑ Loosely inspired by neural connections in the brain
- ❑ Goal: learn hierarchy of features
  - ❑ higher level features are formed by lower level features



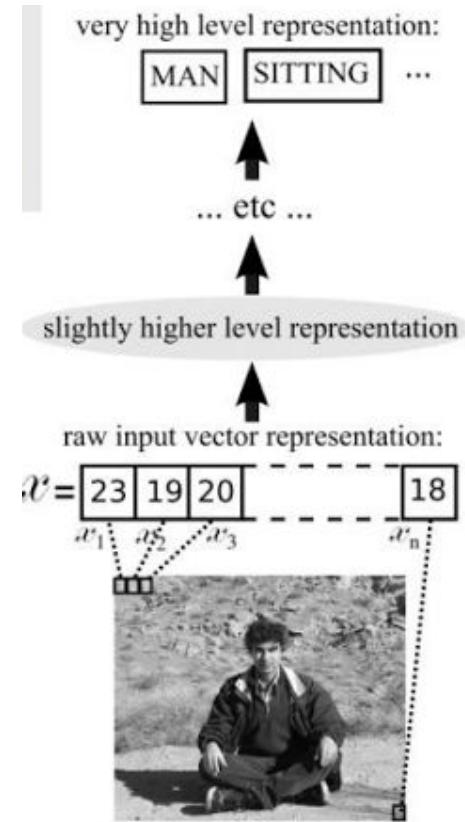
# Deep learning!

- ❑ Loosely inspired by neural connections in the brain
- ❑ Goal: learn hierarchy of features
  - ❑ higher level features are formed by lower level features
- ❑ Used in neuroscience to study intermediate stages of representations



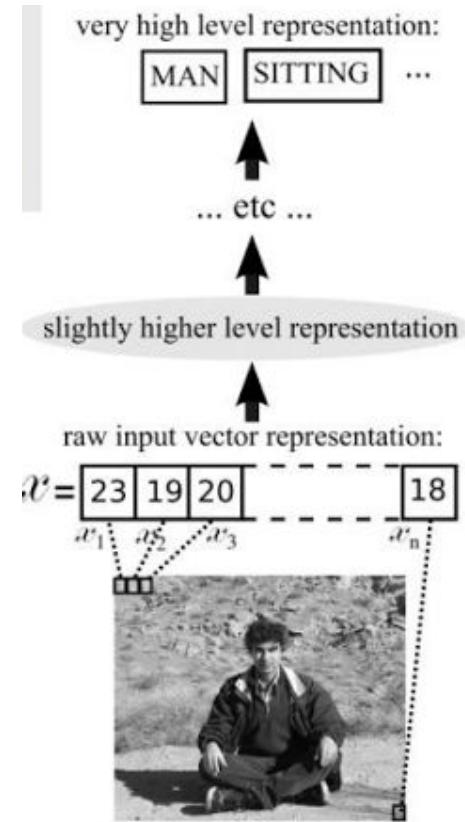
# Deep learning!

- ❑ Loosely inspired by neural connections in the brain
- ❑ Goal: learn hierarchy of features
  - ❑ higher level features are formed by lower level features
- ❑ Used in neuroscience to study intermediate stages of representations
- ❑ There are a few important types:

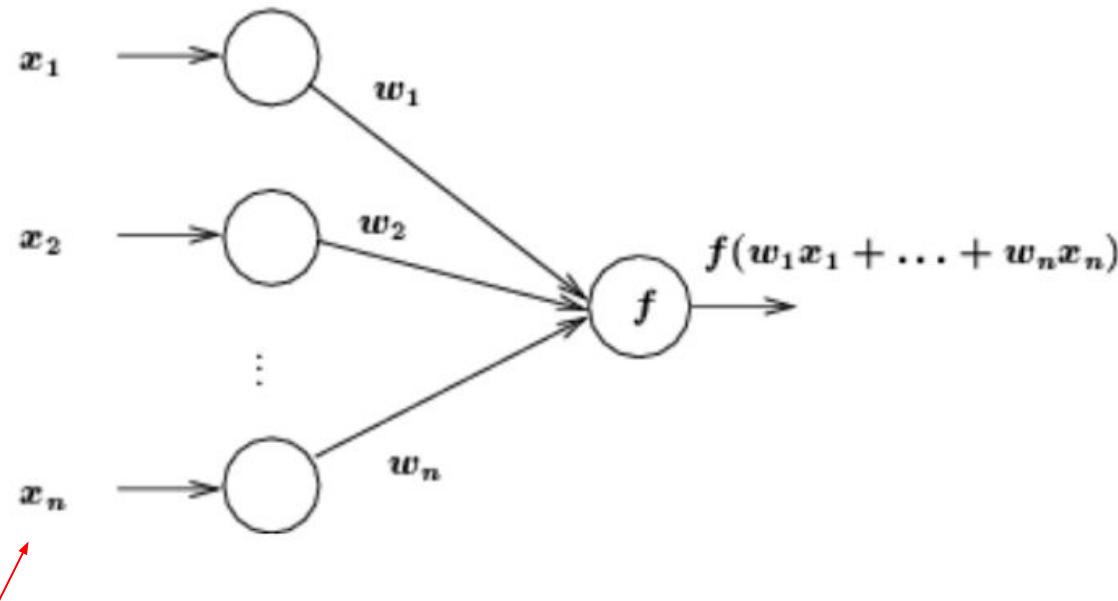


# Deep learning!

- ❑ Loosely inspired by neural connections in the brain
- ❑ Goal: learn hierarchy of features
  - ❑ higher level features are formed by lower level features
- ❑ Used in neuroscience to study intermediate stages of representations
- ❑ There are a few important types:
  - ❑ Neural networks
  - ❑ Recurrent neural networks (RNNs)
    - ❑ Long short-term memory network (LSTM)
  - ❑ Deep belief networks (DBNs)
  - ❑ Convolutional neural networks (CNNs)

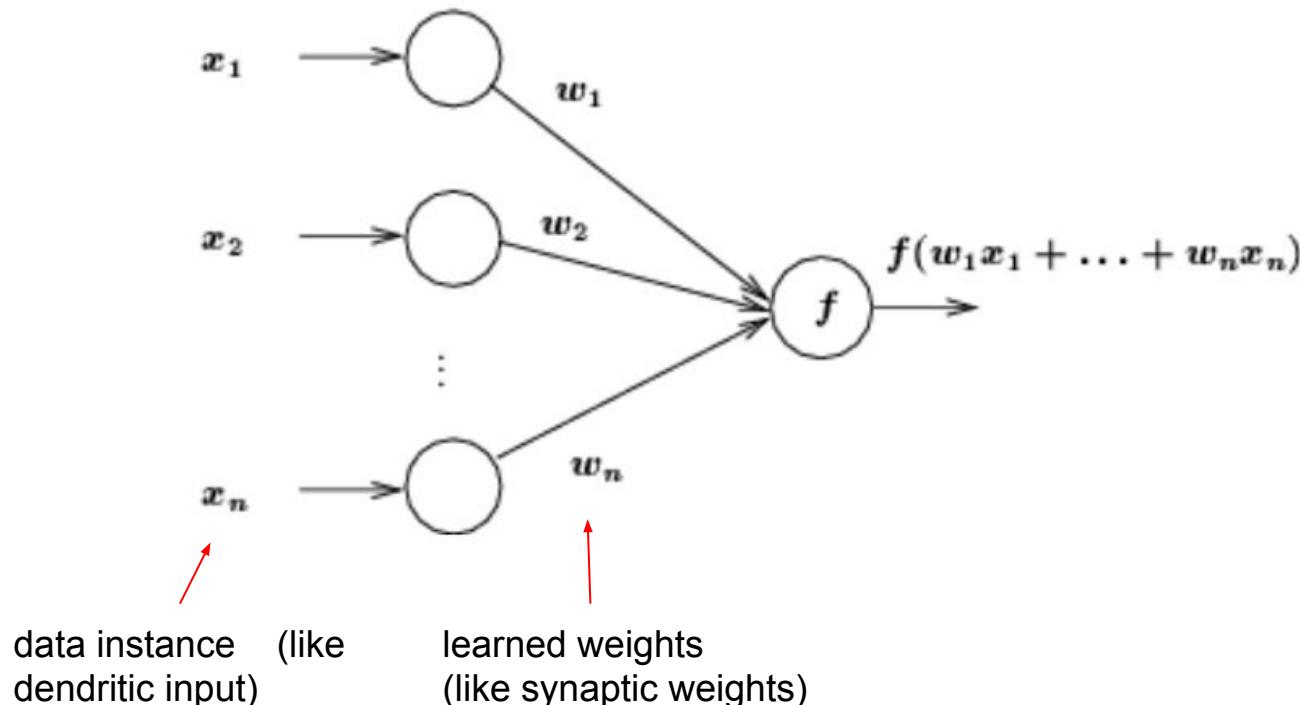


# Artificial neural networks are inspired by connections between real neurons

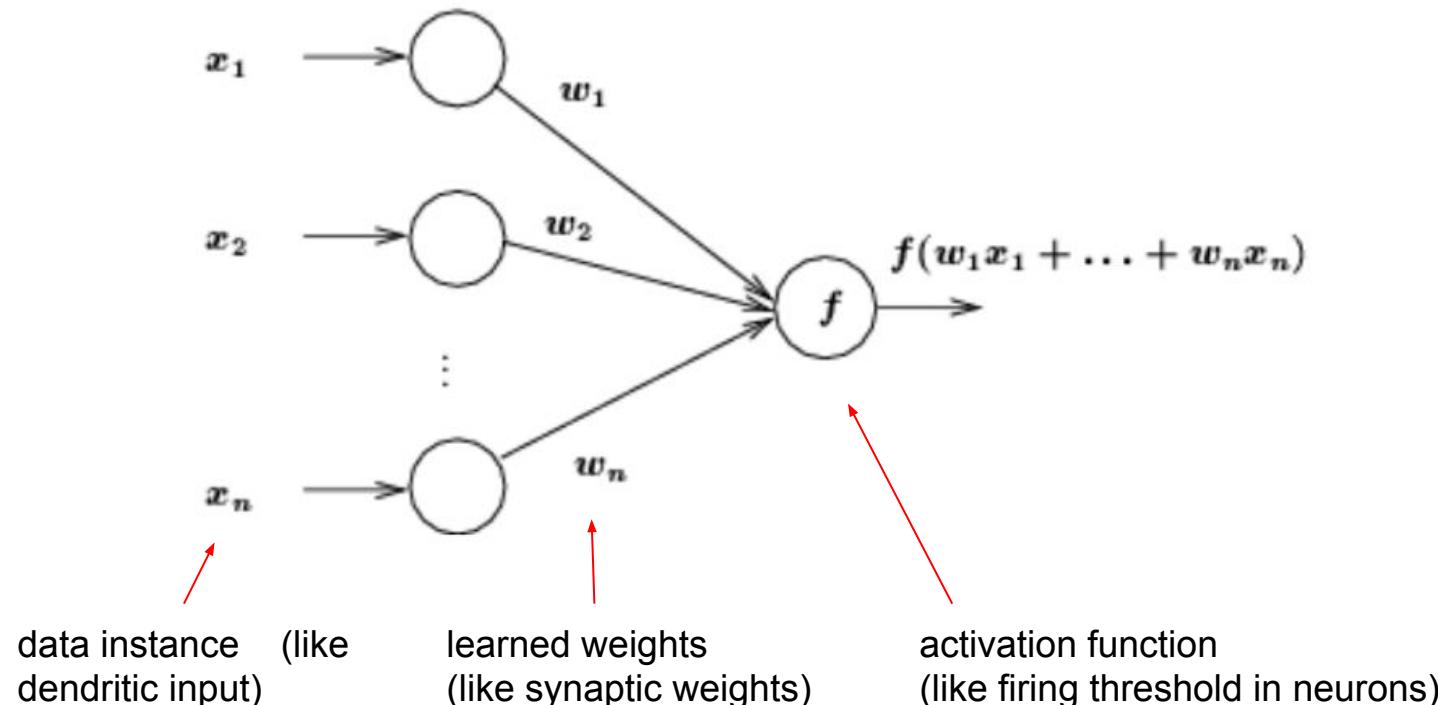


data instance (like  
dendritic input)

# Artificial neural networks are inspired by connections between real neurons

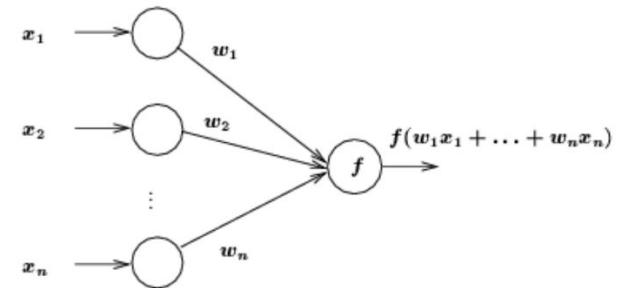


# Artificial neural networks are inspired by connections between real neurons



# What is a good activation function $f$ ?

- Want non-linear  $f$



# What is a good activation function $f$ ?

- Want non-linear  $f$
- Some typical examples:
  - logistic (or sigmoid) function

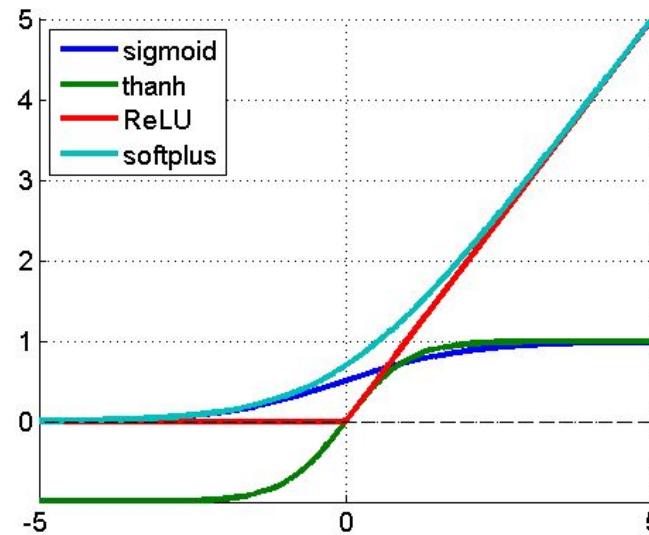
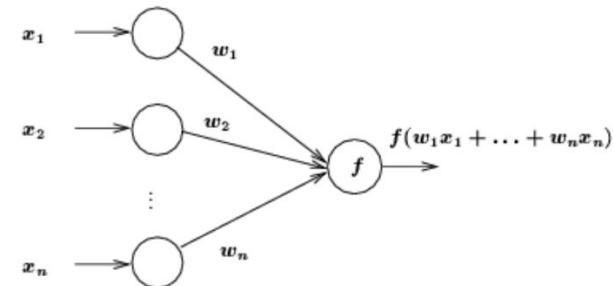
$$f(x) = (1 + e^{-x})^{-1}$$

- hyperbolic tangent function

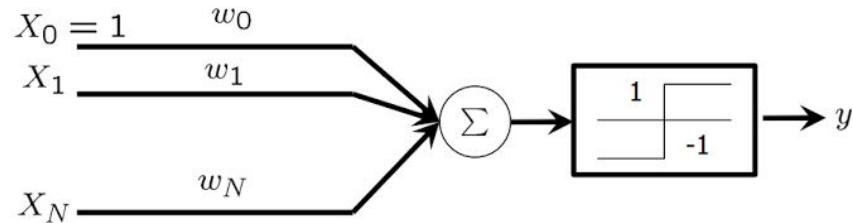
$$f(x) = \tanh(x)$$

- rectified linear function

$$f(x) = \max(0, x)$$



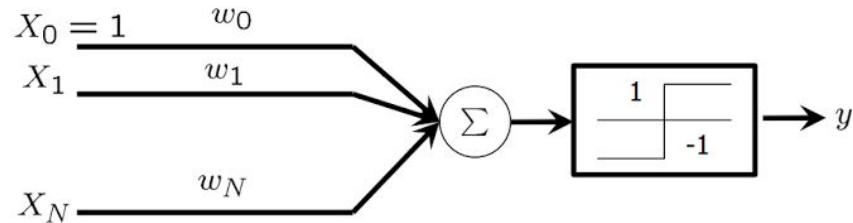
# Perceptron: simplest neural network



$$y = \text{sgn}(\mathbf{w}^T \mathbf{x})$$

- activation function is just a threshold at 0 (checks for the sign of the weighted input sum)

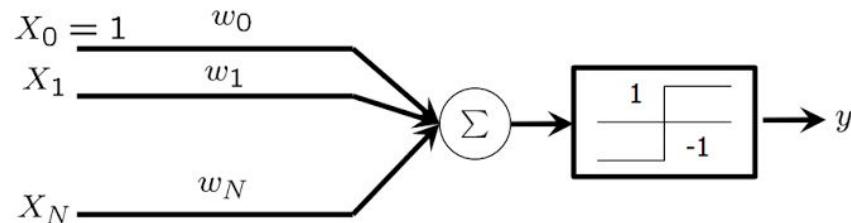
# Perceptron: simplest neural network



$$y = \text{sgn}(\mathbf{w}^T \mathbf{x})$$

- ❑ activation function is just a threshold at 0 (checks for the sign of the weighted input sum)
- ❑ no hidden layer

# Perceptron: simplest neural network

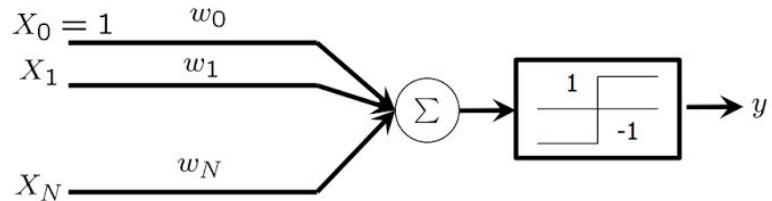


$$y = \text{sgn}(\mathbf{w}^T \mathbf{x})$$

- ❑ activation function is just a threshold at 0 (checks for the sign of the weighted input sum)
- ❑ no hidden layer
- ❑ goal: given the input  $x$  and label  $y$ , find weights  $w$  such that  $y = \text{sgn}(w^T x)$

# How do we learn the weights in the perceptron?

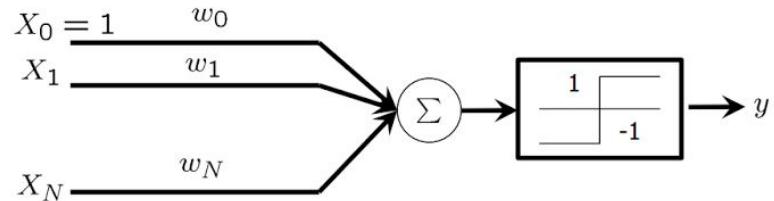
- Key idea: find classification error and update weights to reduce this error



$$y = \text{sgn}(\mathbf{w}^T \mathbf{x})$$

# How do we learn the weights in the perceptron?

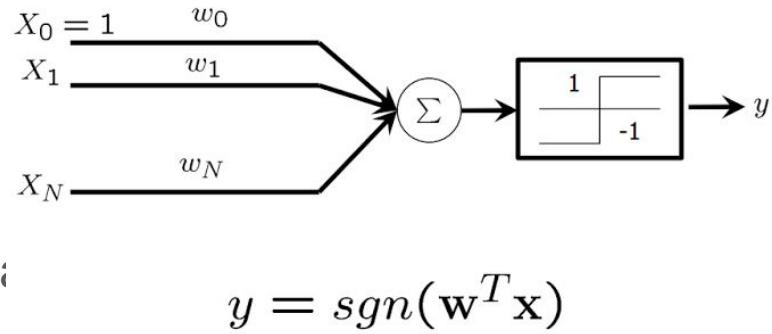
- ❑ Key idea: find classification error and update weights to reduce this error
- ❑ step 1: initialize  $w_0, \dots, w_N$  randomly



$$y = \text{sgn}(\mathbf{w}^T \mathbf{x})$$

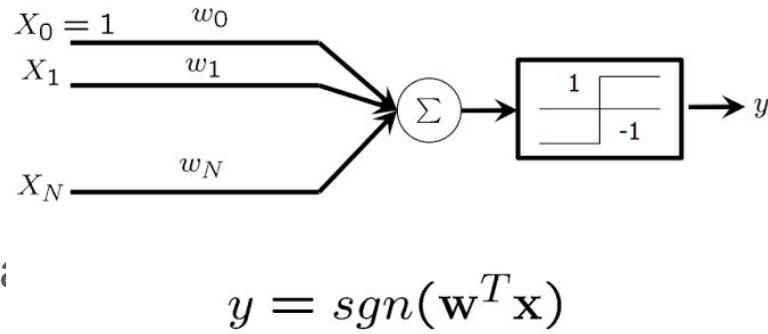
# How do we learn the weights in the perceptron?

- ❑ Key idea: find classification error and update weights to reduce this error
- ❑ step 1: initialize  $w_0, \dots, w_N$  randomly
- ❑ step 2: let  $x^k$  be a training data instance misclassified by  $w(t-1)$



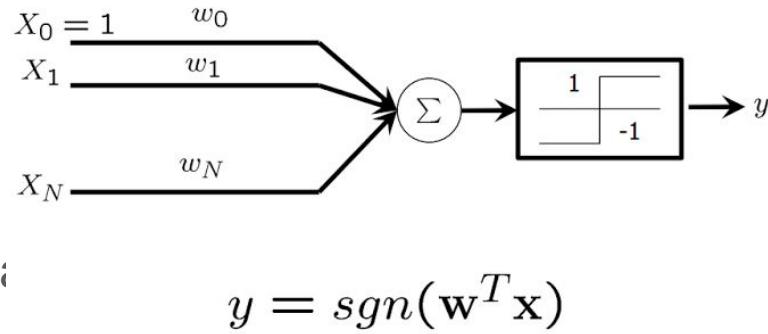
# How do we learn the weights in the perceptron?

- ❑ Key idea: find classification error and update weights to reduce this error
- ❑ step 1: initialize  $w_0, \dots, w_N$  randomly
- ❑ step 2: let  $x^k$  be a training data instance misclassified by  $w(t-1)$
- ❑ step 3: if there is such an instance, calculate classification error  $\epsilon(t) = y^k - \text{sgn}(w(t-1)^T x^k)$



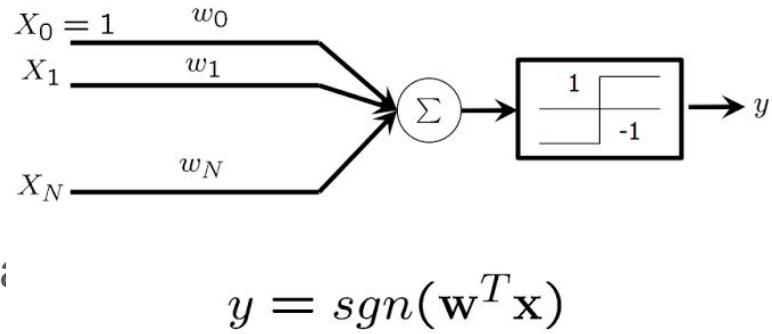
# How do we learn the weights in the perceptron?

- ❑ Key idea: find classification error and update weights to reduce this error
- ❑ step 1: initialize  $w_0, \dots, w_N$  randomly
- ❑ step 2: let  $x^k$  be a training data instance misclassified by  $w(t-1)$
- ❑ step 3: if there is such an instance, calculate classification error  $\varepsilon(t) = y^k - \text{sgn}(w(t-1)^T x^k)$
- ❑ step 4: account for error by updating weights:
  - ❑  $w(t) = w(t-1) + \alpha \varepsilon(t) x^k$



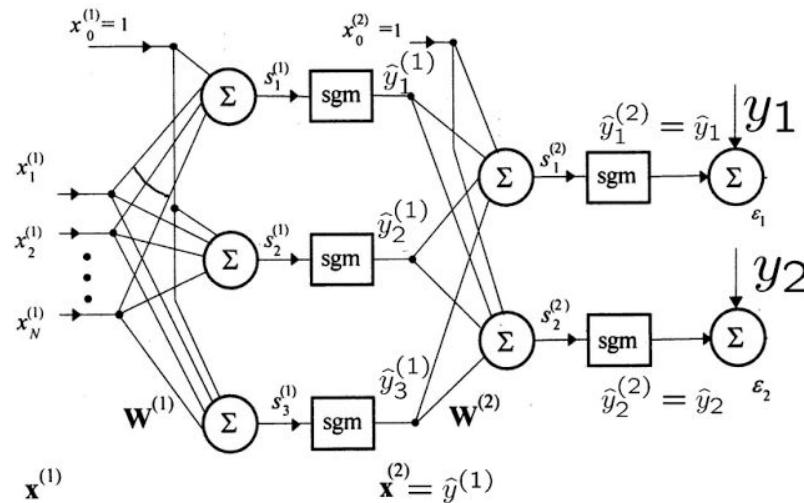
# How do we learn the weights in the perceptron?

- ❑ Key idea: find classification error and update weights to reduce this error
- ❑ step 1: initialize  $w_0, \dots, w_N$  randomly
- ❑ step 2: let  $x^k$  be a training data instance misclassified by  $w(t-1)$
- ❑ step 3: if there is such an instance, calculate classification error  $\epsilon(t) = y^k - \text{sgn}(w(t-1)^T x^k)$
- ❑ step 4: account for error by updating weights:
  - ❑  $w(t) = w(t-1) + \alpha \epsilon(t) x^k$
- ❑ step 5:  $t = t+1$ , end when there are no misclassifications



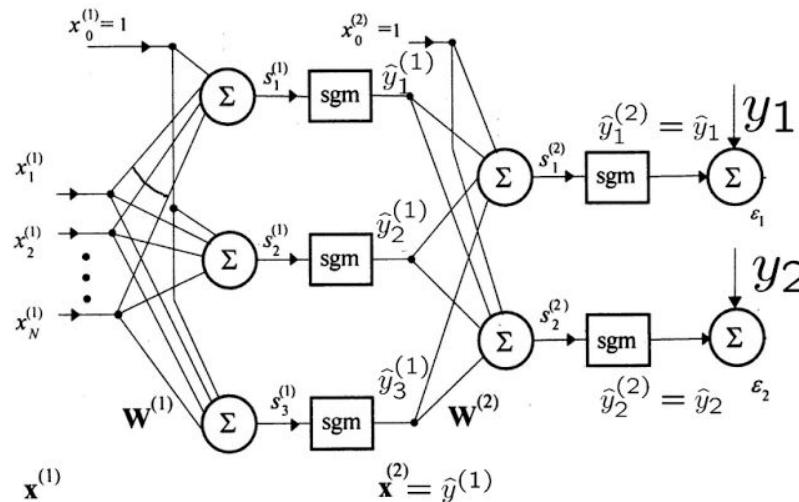
# We add more layers => multilayer perceptron (MLP)

- Additional layers (not input or output) are called “hidden” layers



# We add more layers => multilayer perceptron (MLP)

- Additional layers (not input or output) are called “hidden” layers



- Still need to learn weights  $W$ , but now we have multiple layers

# Learning weights with multiple layers: backpropagation of error

- When there is a misclassification, different neurons have different amount of responsibility for it

# Learning weights with multiple layers: backpropagation of error

- ❑ When there is a misclassification, different neurons have different amount of responsibility for it
- ❑ Key observation: the misclassification error for a neuron in layer  $m$  can be calculated from the misclassification errors for all of the neurons in layer  $m+1$ 
  - ❑ because the network is fully connected

# Learning weights with multiple layers: backpropagation of error

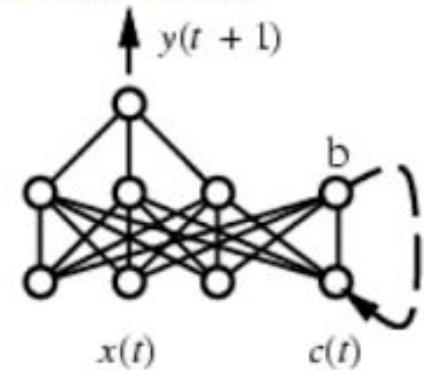
- ❑ When there is a misclassification, different neurons have different amount of responsibility for it
- ❑ Key observation: the misclassification error for a neuron in layer  $m$  can be calculated from the misclassification errors for all of the neurons in layer  $m+1$ 
  - ❑ because the network is fully connected
- ❑ So we can calculate the final misclassification error and propagate it backward

# Learning weights with multiple layers: backpropagation of error

- ❑ When there is a misclassification, different neurons have different amount of responsibility for it
- ❑ Key observation: the misclassification error for a neuron in layer  $m$  can be calculated from the misclassification errors for all of the neurons in layer  $m+1$ 
  - ❑ because the network is fully connected
- ❑ So we can calculate the final misclassification error and propagate it backward
- ❑ Update weights according to the error

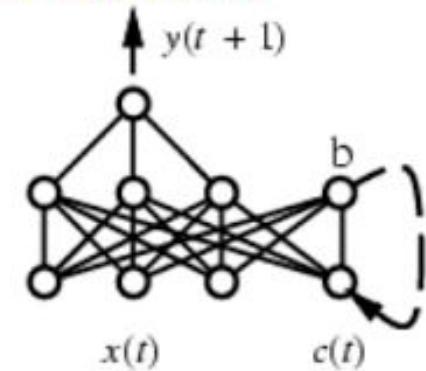
# Neural networks for sequential data: recurrent NNs

- In sequential data, history can help predict the future



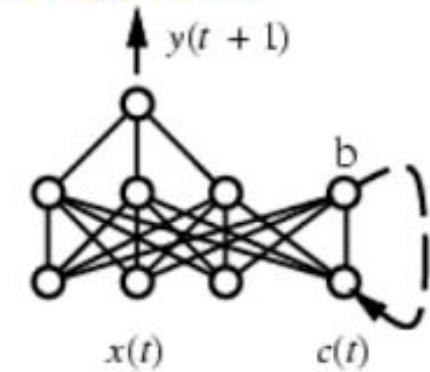
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network



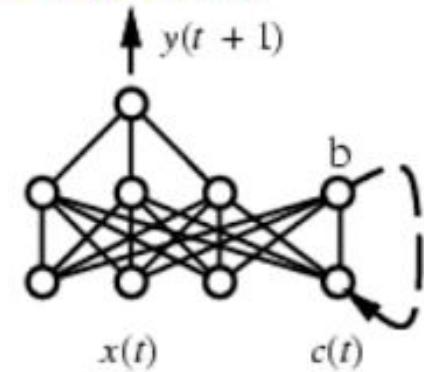
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state



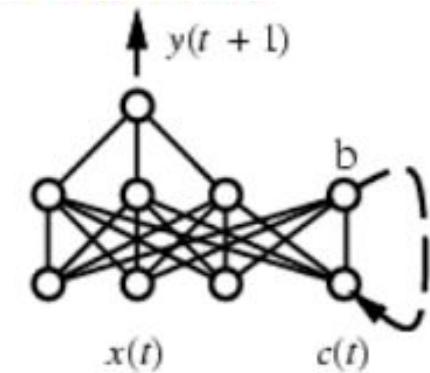
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state
- ❑ In regular RNNs, the recurrent state is overwritten frequently



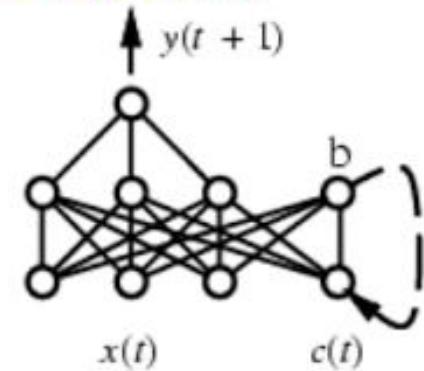
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state
- ❑ In regular RNNs, the recurrent state is overwritten frequently
  - ❑ Good if important events happen close together



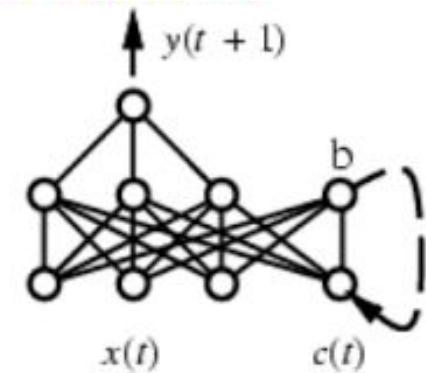
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state
- ❑ In regular RNNs, the recurrent state is overwritten frequently
  - ❑ Good if important events happen close together
- ❑ But what if there are large delays before important events?



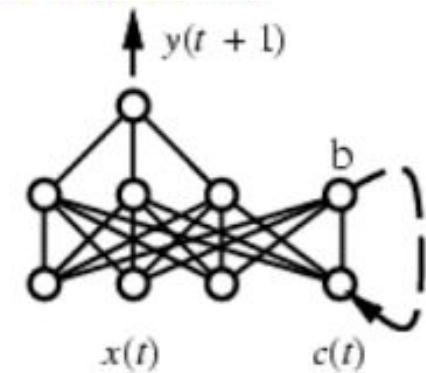
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state
- ❑ In regular RNNs, the recurrent state is overwritten frequently
  - ❑ Good if important events happen close together
- ❑ But what if there are large delays before important events?
  - ❑ Long short-term memory (LSTM) networks!



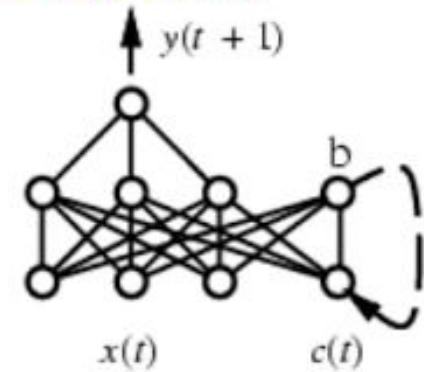
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state
- ❑ In regular RNNs, the recurrent state is overwritten frequently
  - ❑ Good if important events happen close together
- ❑ But what if there are large delays before important events?
  - ❑ Long short-term memory (LSTM) networks!
    - ❑ Best known result in connected handwriting recognition



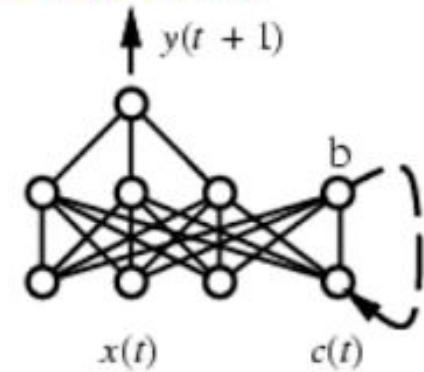
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state
- ❑ In regular RNNs, the recurrent state is overwritten frequently
  - ❑ Good if important events happen close together
- ❑ But what if there are large delays before important events?
  - ❑ Long short-term memory (LSTM) networks!
    - ❑ Best known result in connected handwriting recognition
    - ❑ Also used for automatic speech recognition



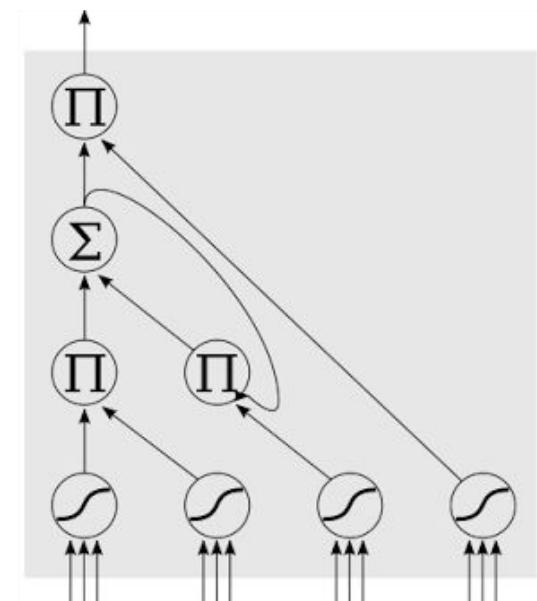
# Neural networks for sequential data: recurrent NNs

- ❑ In sequential data, history can help predict the future
- ❑ Would like to explicitly capture state history in the neural network
  - ❑ Add a recurrent state
- ❑ In regular RNNs, the recurrent state is overwritten frequently
  - ❑ Good if important events happen close together
- ❑ But what if there are large delays before important events?
  - ❑ Long short-term memory (LSTM) networks!
    - ❑ Best known result in connected handwriting recognition
    - ❑ Also used for automatic speech recognition
    - ❑ In 2016, Google, Apple, Microsoft, Baidu reveal LSTMs as fundamental components in their technologies



# LSTMs: remembering values for arbitrary length of time

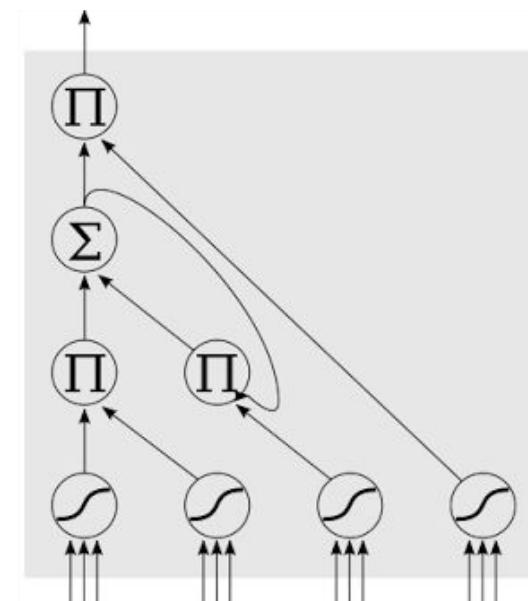
- ❑ Contains LSTM blocks



input value	input gate	forget gate	output gate
----------------	---------------	----------------	----------------

# LSTMs: remembering values for arbitrary length of time

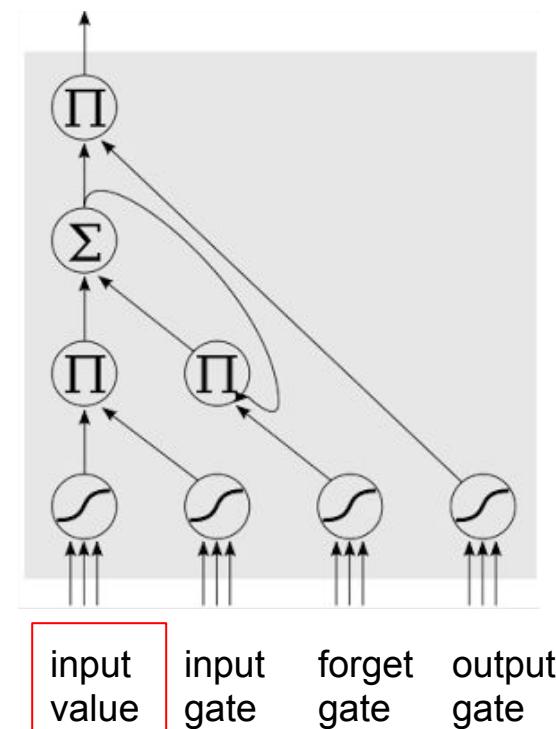
- ❑ Contains LSTM blocks
  - ❑ each block contains gates that determine when the input is significant enough to remember



input value	input gate	forget gate	output gate
-------------	------------	-------------	-------------

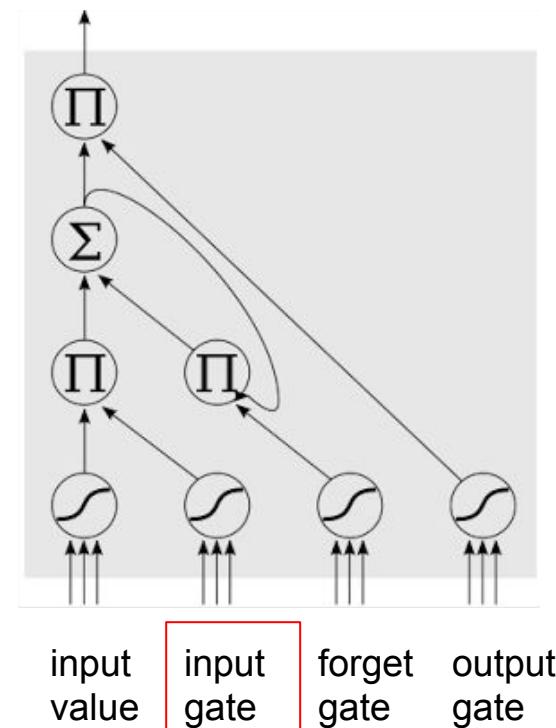
# LSTMs: remembering values for arbitrary length of time

- ❑ Contains LSTM blocks
  - ❑ each block contains gates that determine when the input is significant enough to remember



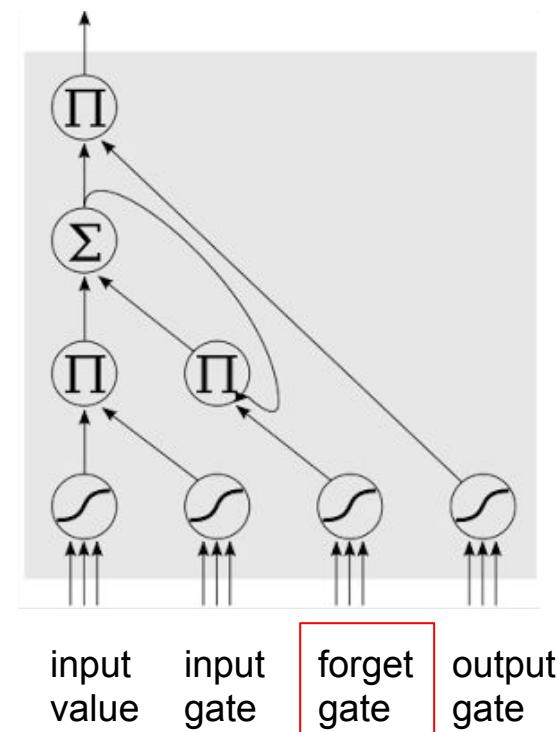
# LSTMs: remembering values for arbitrary length of time

- ❑ Contains LSTM blocks
  - ❑ each block contains gates that determine when the input is significant enough to remember



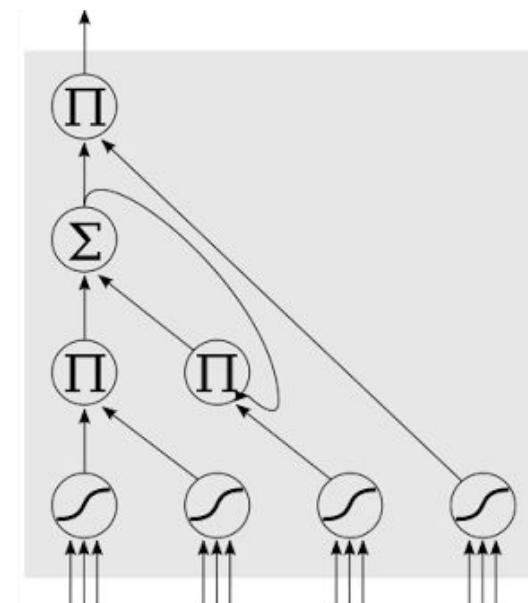
# LSTMs: remembering values for arbitrary length of time

- ❑ Contains LSTM blocks
  - ❑ each block contains gates that determine when the input is significant enough to remember



# LSTMs: remembering values for arbitrary length of time

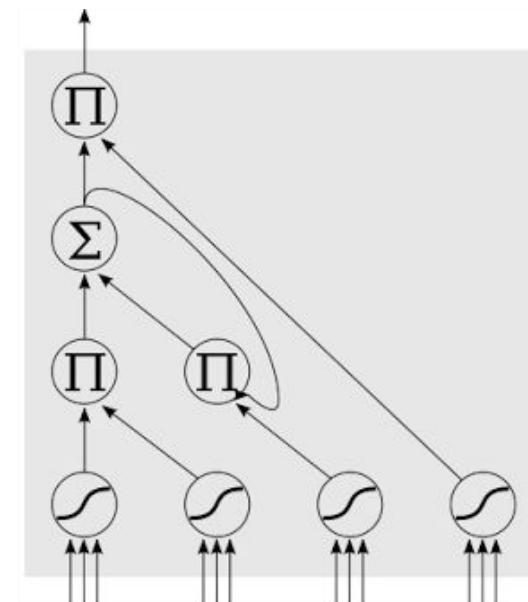
- ❑ Contains LSTM blocks
  - ❑ each block contains gates that determine when the input is significant enough to remember



input value	input gate	forget gate	output gate
-------------	------------	-------------	-------------

# LSTMs: remembering values for arbitrary length of time

- ❑ Contains LSTM blocks
  - ❑ each block contains gates that determine when the input is significant enough to remember
- ❑ LSTMs, like regular RNNs, can be trained with backpropagation



input value	input gate	forget gate	output gate
----------------	---------------	----------------	----------------

# 3 main problems with backpropagation

- ❑ Problem 1: random initialization of weights can lead to bad local minimum

# 3 main problems with backpropagation

- ❑ Problem 1: random initialization of weights can lead to bad local minimum
- ❑ Problem 2: learning time does not scale well with number of hidden layers => slow

# 3 main problems with backpropagation

- ❑ Problem 1: random initialization of weights can lead to bad local minimum
- ❑ Problem 2: learning time does not scale well with number of hidden layers => slow
- ❑ Problem 3: need a lot of data for good estimates of weights when we start from random initializations and to calculate classification error, we need labels => need a lot of labeled data, which is often scarce

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task
  - ❑ avoids random initialization (problem 1 of backpropagation)

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task
    - ❑ avoids random initialization (problem 1 of backpropagation)
  - ❑ reduces need for a lot of labeled data (problem 3)

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task
    - ❑ avoids random initialization (problem 1 of backpropagation)
    - ❑ reduces need for a lot of labeled data (problem 3)
- ❑ Once the DBN is pre-trained, it is fine-tuned with backpropagation for a particular task

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task
    - ❑ avoids random initialization (problem 1 of backpropagation)
  - ❑ reduces need for a lot of labeled data (problem 3)
- ❑ Once the DBN is pre-trained, it is fine-tuned with backpropagation for a particular task
  - ❑ reduces the learning time of backprop (problem 2)

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task
    - ❑ avoids random initialization (problem 1 of backpropagation)
    - ❑ reduces need for a lot of labeled data (problem 3)
- ❑ Once the DBN is pre-trained, it is fine-tuned with backpropagation for a particular task
  - ❑ reduces the learning time of backprop (problem 2)
- ❑ So why is it only a partial solution?

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task
    - ❑ avoids random initialization (problem 1 of backpropagation)
    - ❑ reduces need for a lot of labeled data (problem 3)
- ❑ Once the DBN is pre-trained, it is fine-tuned with backpropagation for a particular task
  - ❑ reduces the learning time of backprop (problem 2)
- ❑ So why is it only a partial solution?
  - ❑ pre-training is done in an unsupervised way, in which the objective is to minimize some reconstruction error

# Partial solution: deep belief networks (DBN)

- ❑ DBNs use a technique called “pre-training” to train each pair of successive layers separately without using labeled data and backpropagation
  - ❑ uses the resulting pre-trained weights to initialize the DBN weights for the task
    - ❑ avoids random initialization (problem 1 of backpropagation)
  - ❑ reduces need for a lot of labeled data (problem 3)
- ❑ Once the DBN is pre-trained, it is fine-tuned with backpropagation for a particular task
  - ❑ reduces the learning time of backprop (problem 2)
- ❑ So why is it only a partial solution?
  - ❑ pre-training is done in an unsupervised way, in which the objective is to minimize some reconstruction error
  - ❑ in practice, many times weights learned in this way are not close enough to the ones needed for a classification task, and bad local minima are not avoided

# Another partial solution: convolutional neural networks (CNNs)

- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)

# Another partial solution: convolutional neural networks (CNNs)

- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)
- ❑ Sparse connections between layers (not fully connected like regular NNs)
  - ❑ reduces learning time (problem 2)

# Another partial solution: convolutional neural networks (CNNs)

- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)
- ❑ Sparse connections between layers (not fully connected like regular NNs)
  - ❑ reduces learning time (problem 2)
- ❑ Often, researchers take an already trained CNN on a related task and fine-tune it for the specific task

# Another partial solution: convolutional neural networks (CNNs)

- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)
- ❑ Sparse connections between layers (not fully connected like regular NNs)
  - ❑ reduces learning time (problem 2)
- ❑ Often, researchers take an already trained CNN on a related task and fine-tune it for the specific task
  - ❑ avoids random initialization of weights (problem 1)

# Another partial solution: convolutional neural networks (CNNs)

- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)
- ❑ Sparse connections between layers (not fully connected like regular NNs)
  - ❑ reduces learning time (problem 2)
- ❑ Often, researchers take an already trained CNN on a related task and fine-tune it for the specific task
  - ❑ avoids random initialization of weights (problem 1)
  - ❑ reduces need for labeled data for learning (problem 3)

# Another partial solution: convolutional neural networks (CNNs)

- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)
- ❑ Sparse connections between layers (not fully connected like regular NNs)
  - ❑ reduces learning time (problem 2)
- ❑ Often, researchers take an already trained CNN on a related task and fine-tune it for the specific task
  - ❑ avoids random initialization of weights (problem 1)
  - ❑ reduces need for labeled data for learning (problem 3)
- ❑ So why only partial solution?

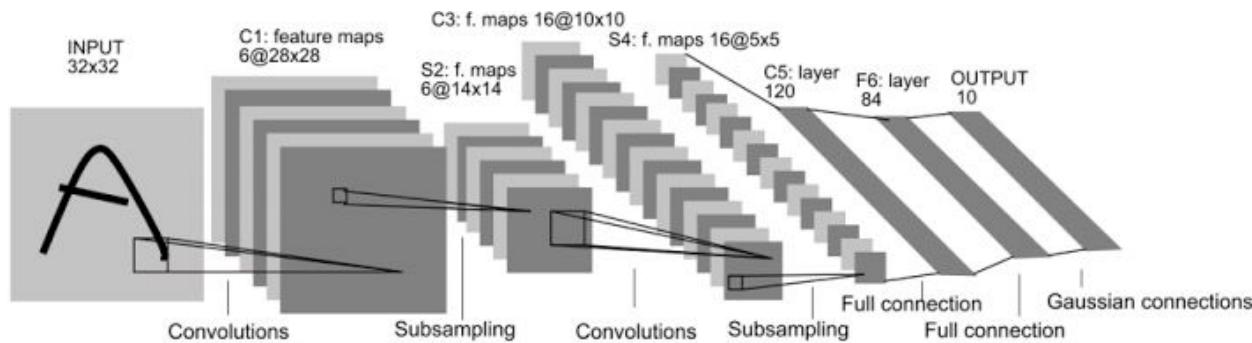
# Another partial solution: convolutional neural networks (CNNs)

- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)
- ❑ Sparse connections between layers (not fully connected like regular NNs)
  - ❑ reduces learning time (problem 2)
- ❑ Often, researchers take an already trained CNN on a related task and fine-tune it for the specific task
  - ❑ avoids random initialization of weights (problem 1)
  - ❑ reduces need for labeled data for learning (problem 3)
- ❑ So why only partial solution?
  - ❑ a CNN trained on a related task is not always available => back to problems 1 and 3

# Another partial solution: convolutional neural networks (CNNs)

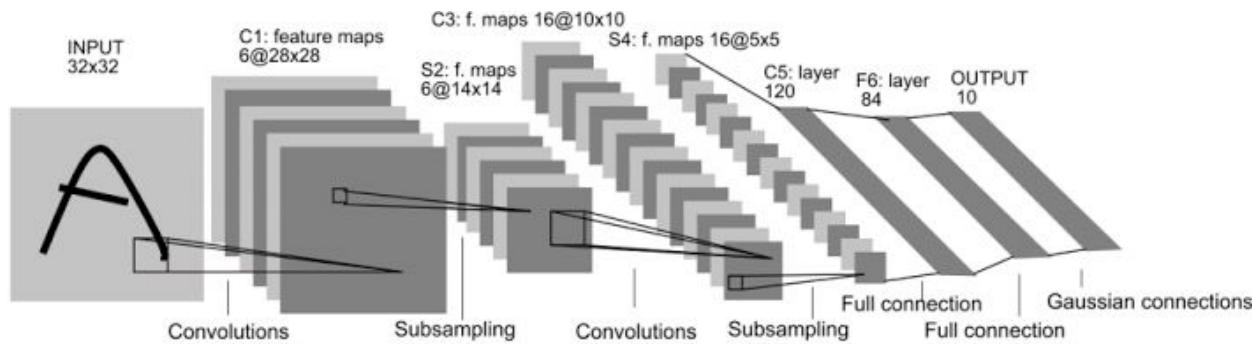
- ❑ Inspired by the structure of the visual system (receptive fields of simple and complex cells)
- ❑ Sparse connections between layers (not fully connected like regular NNs)
  - ❑ reduces learning time (problem 2)
- ❑ Often, researchers take an already trained CNN on a related task and fine-tune it for the specific task
  - ❑ avoids random initialization of weights (problem 1)
  - ❑ reduces need for labeled data for learning (problem 3)
- ❑ So why only partial solution?
  - ❑ a CNN trained on a related task is not always available => back to problems 1 and 3
  - ❑ though, for vision tasks (e.g. classifying images), a good CNN is available

# CNN: alternating convolutional and downsampling layers



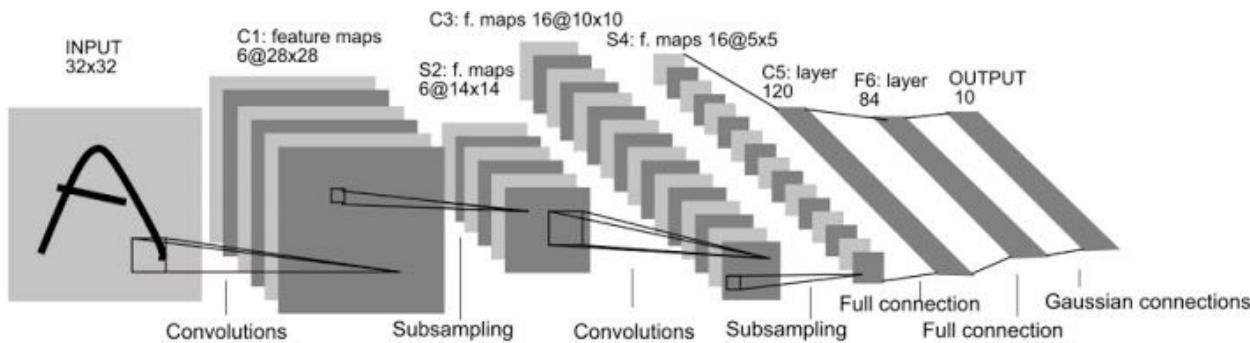
- Restricted connections => makes learning the weights more tractable

# CNN: alternating convolutional and downsampling layers



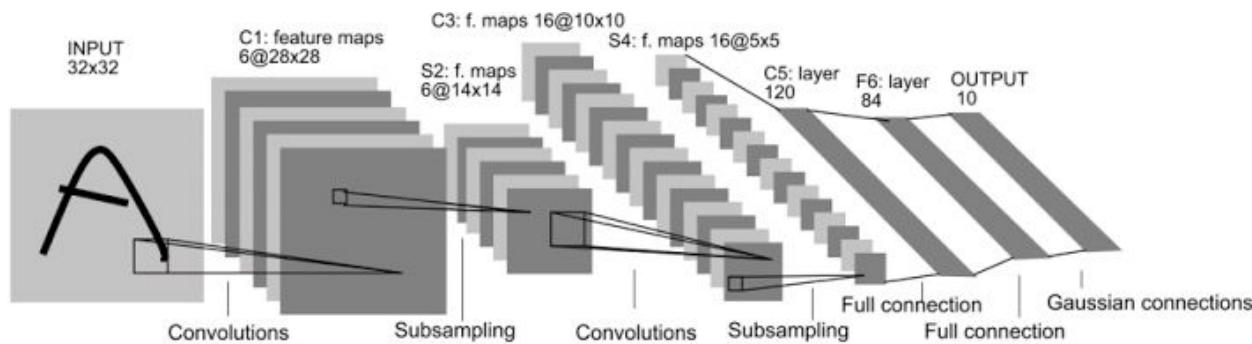
- ❑ Restricted connections => makes learning the weights more tractable
- ❑ Convolutional layers = each neuron applies some kernel to its receptive field

# CNN: alternating convolutional and downsampling layers



- ❑ Restricted connections => makes learning the weights more tractable
- ❑ Convolutional layers = each neuron applies some kernel to its receptive field
- ❑ Subsampling layers = reduce dimensionality, maintain translational invariance

# CNN: alternating convolutional and downsampling layers



- ❑ Restricted connections => makes learning the weights more tractable
- ❑ Convolutional layers = each neuron applies some kernel to its receptive field
- ❑ Subsampling layers = reduce dimensionality, maintain translational invariance
- ❑ Shared weights between feature maps reduce number of parameters to learn

# Most frequently used CNN is trained on ImageNet

- ImageNet is a database of 15M images collected from the web



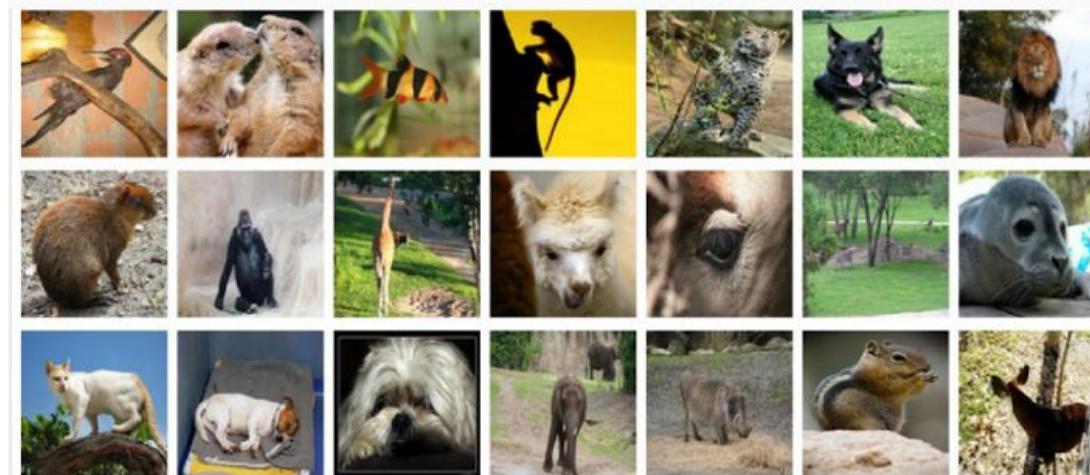
# Most frequently used CNN is trained on ImageNet

- ImageNet is a database of 15M images collected from the web
- 22k categories



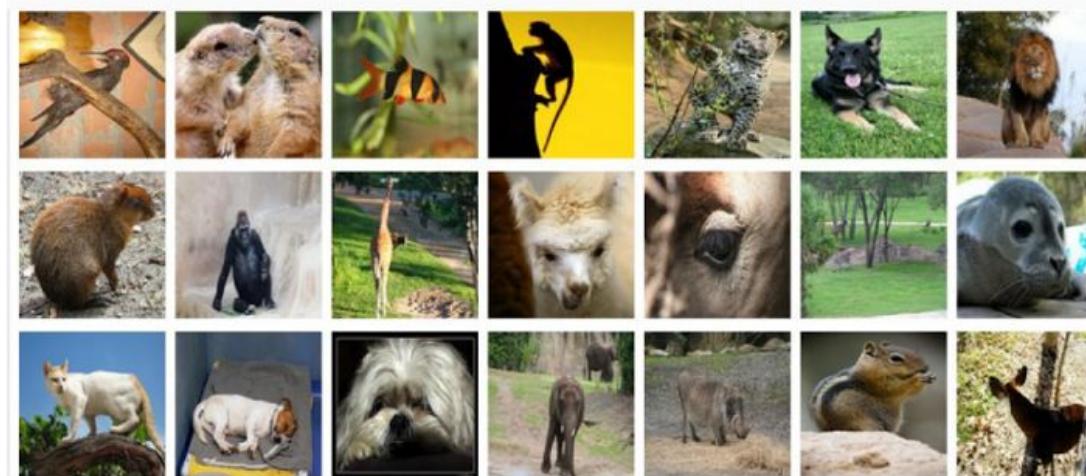
# Most frequently used CNN is trained on ImageNet

- ImageNet is a database of 15M images collected from the web
- 22k categories
- labels provided by people on MTurk



# Most frequently used CNN is trained on ImageNet

- ImageNet is a database of 15M images collected from the web
- 22k categories
- labels provided by people on MTurk
- RGB images (not black and white)



# CNN results in ImageNet classification challenge

- ❑ 1k categories, 1.2M training images, 50k validation, 150k testing

# CNN results in ImageNet classification challenge

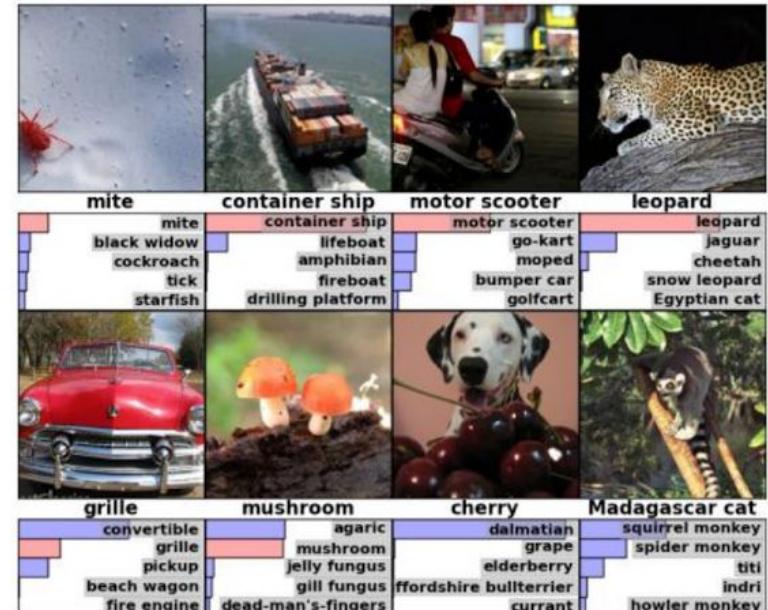
- ❑ 1k categories, 1.2M training images, 50k validation, 150k testing
- ❑ Classification challenge:
  - ❑ Make 1 guess about the label (top-1 error)
  - ❑ Make 5 guesses about the label (top-5 error)

# CNN results in ImageNet classification challenge

- ❑ 1k categories, 1.2M training images, 50k validation, 150k testing
- ❑ Classification challenge:
  - ❑ Make 1 guess about the label (top-1 error)
  - ❑ Make 5 guesses about the label (top-5 error)
- ❑ CNN results (won)
  - ❑ Top-1 error: 37.5%
  - ❑ Top-5 error: 17.0%

# CNN results in ImageNet classification challenge

- ❑ 1k categories, 1.2M training images, 50k validation, 150k testing
- ❑ Classification challenge:
  - ❑ Make 1 guess about the label (top-1 error)
  - ❑ Make 5 guesses about the label (top-5 error)
- ❑ CNN results (won)
  - ❑ Top-1 error: 37.5%
  - ❑ Top-5 error: 17.0%



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal
  - ❑ Total number of possible board positions =  $3^{19 \times 19} \times 1.2\% = 2.1 \times 10^{170} > \# \text{ atoms in universe}$



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal
    - ❑ Total number of possible board positions =  $3^{19 \times 19} \times 1.2\% = 2.1 \times 10^{170} > \# \text{ atoms in universe}$
- ❑ So enumerating all possible board outcomes and scoring them won't work



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal
    - ❑ Total number of possible board positions =  $3^{19 \times 19} \times 1.2\% = 2.1 \times 10^{170} > \# \text{ atoms in universe}$
- ❑ So enumerating all possible board outcomes and scoring them won't work
- ❑ Best approach before AlphaGo:



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal
    - ❑ Total number of possible board positions =  $3^{19 \times 19} \times 1.2\% = 2.1 \times 10^{170} > \# \text{ atoms in universe}$
- ❑ So enumerating all possible board outcomes and scoring them won't work
- ❑ Best approach before AlphaGo:
  - ❑ Tree search enhanced by policies that were trained to predict human expert's moves



# CNN also used in AlphaGo

- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal
    - ❑ Total number of possible board positions =  $3^{19 \times 19} \times 1.2\% = 2.1 \times 10^{170} > \# \text{ atoms in universe}$
- ❑ So enumerating all possible board outcomes and scoring them won't work
- ❑ Best approach before AlphaGo:
  - ❑ Tree search enhanced by policies that were trained to predict human expert's moves
  - ❑ "Weak amateur level play"



# CNN also used in AlphaGo



- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal
    - ❑ Total number of possible board positions =  $3^{19 \times 19} \times 1.2\% = 2.1 \times 10^{170} > \# \text{ atoms in universe}$
- ❑ So enumerating all possible board outcomes and scoring them won't work
- ❑ Best approach before AlphaGo:
  - ❑ Tree search enhanced by policies that were trained to predict human expert's moves
  - ❑ "Weak amateur level play"
- ❑ AlphaGo combines tree search with RL executed with policy and value CNN networks

# CNN also used in AlphaGo



- ❑ First off, why is Go so hard?
  - ❑ larger board than chess (19x19)
  - ❑ many more possible board outcomes than chess
    - ❑ each board position has 3 possibilities: empty, black, white
    - ❑ an estimated 1.2% of all board positions are legal
    - ❑ Total number of possible board positions =  $3^{19 \times 19} \times 1.2\% = 2.1 \times 10^{170} > \# \text{ atoms in universe}$
- ❑ So enumerating all possible board outcomes and scoring them won't work
- ❑ Best approach before AlphaGo:
  - ❑ Tree search enhanced by policies that were trained to predict human expert's moves
  - ❑ "Weak amateur level play"
- ❑ AlphaGo combines tree search with RL executed with policy and value CNN networks
  - ❑ Won 4 out of 5 games against Lee Sedol, second best player of Go in the world (by # of won

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value
- ❑ Input to CNNs: board position (19x19 image)

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value
- ❑ Input to CNNs: board position (19x19 image)
- ❑ First policy CNN trained to select actions that predict human moves selected in a given state

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value
- ❑ Input to CNNs: board position (19x19 image)
- ❑ First policy CNN trained to select actions that predict human moves selected in a given state
  - ❑ previous methods used shallow classification methods to do this

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value
- ❑ Input to CNNs: board position (19x19 image)
- ❑ First policy CNN trained to select actions that predict human moves selected in a given state
  - ❑ previous methods used shallow classification methods to do this
- ❑ Second policy CNN trained to select actions that maximize expected future reward (winning) on games between AlphaGo's current policy strategy and a randomly selected previous iteration of the policy strategy

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value
- ❑ Input to CNNs: board position (19x19 image)
- ❑ First policy CNN trained to select actions that predict human moves selected in a given state
  - ❑ previous methods used shallow classification methods to do this
- ❑ Second policy CNN trained to select actions that maximize expected future reward (winning) on games between AlphaGo's current policy strategy and a randomly selected previous iteration of the policy strategy
  - ❑ goal: to adjust the policy towards the correct goal of winning games, rather than predictive accuracy

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value
- ❑ Input to CNNs: board position (19x19 image)
- ❑ First policy CNN trained to select actions that predict human moves selected in a given state
  - ❑ previous methods used shallow classification methods to do this
- ❑ Second policy CNN trained to select actions that maximize expected future reward (winning) on games between AlphaGo's current policy strategy and a randomly selected previous iteration of the policy strategy
  - ❑ goal: to adjust the policy towards the correct goal of winning games, rather than predictive accuracy
- ❑ Value CNN trained to predict the winner of games played by the RL policy network against itself

# AlphaGo algorithm details

- ❑ 3 CNNs - 2 for policy, 1 for value
- ❑ Input to CNNs: board position (19x19 image)
- ❑ First policy CNN trained to select actions that predict human moves selected in a given state
  - ❑ previous methods used shallow classification methods to do this
- ❑ Second policy CNN trained to select actions that maximize expected future reward (winning) on games between AlphaGo's current policy strategy and a randomly selected previous iteration of the policy strategy
  - ❑ goal: to adjust the policy towards the correct goal of winning games, rather than predictive accuracy
- ❑ Value CNN trained to predict the winner of games played by the RL policy network against itself
- ❑ Finally, use the value and policy networks to reduce depth breadth of search

# Getting started with deep learning

- ❑ A bit of an art

# Getting started with deep learning

- ❑ A bit of an art
- ❑ Several deep learning libraries

# Getting started with deep learning

- ❑ A bit of an art
- ❑ Several deep learning libraries
  - ❑ Theano
    - ❑ Written in python, compatible with GPUs

# Getting started with deep learning

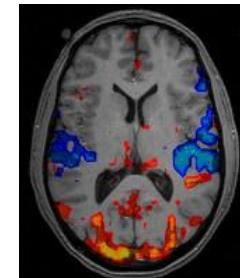
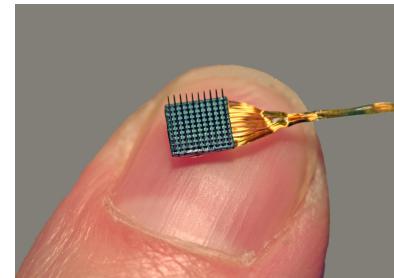
- ❑ A bit of an art
- ❑ Several deep learning libraries
  - ❑ Theano
    - ❑ Written in python, compatible with GPUs
  - ❑ TensorFlow
    - ❑ Similar to Theano but arguably more intuitive
    - ❑ <https://www.tensorflow.org/>

# Main takeaways

- ❑ Considering the sequential nature of time series data is important
- ❑ Ways to account for this time dependence is through models, such as hidden Markov models, Markov decision processes, and long short-term memory neural networks
- ❑ We don't know why exactly deep learning works *yet*, but it's increasingly more popular, both in industry and academia
  - ❑ deep learning course at CMU taught by Ruslan Salakhutdinov in the Fall!

# How can ML help neuroscientists?

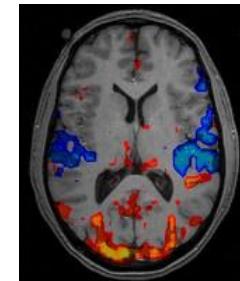
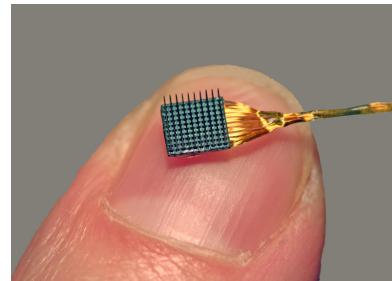
- ❑ Deal with large number of sensors/recording sites



- ❑ investigate high-dimensional representations

# How can ML help neuroscientists?

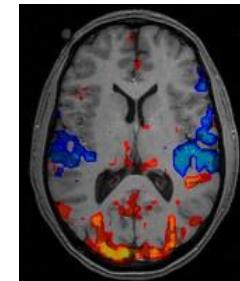
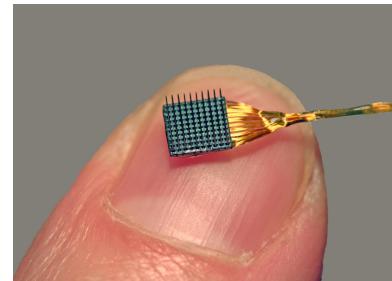
- ❑ Deal with large number of sensors/recording sites



- ❑ investigate high-dimensional representations
  - ❑ classification (what does this high-dimensional data represent?)

# How can ML help neuroscientists?

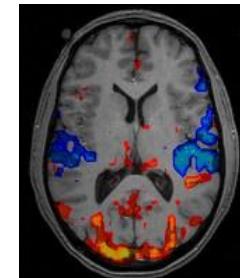
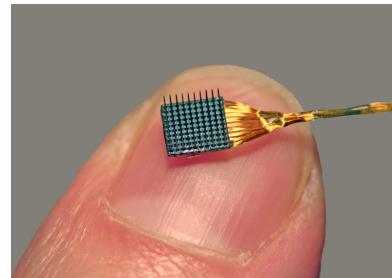
- ❑ Deal with large number of sensors/recording sites



- ❑ investigate high-dimensional representations
  - ❑ classification (what does this high-dimensional data represent?)
  - ❑ regression (how does it represent it? can we predict a different representation?)

# How can ML help neuroscientists?

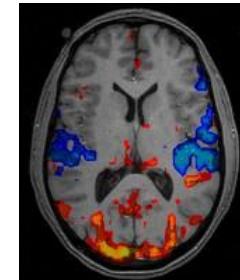
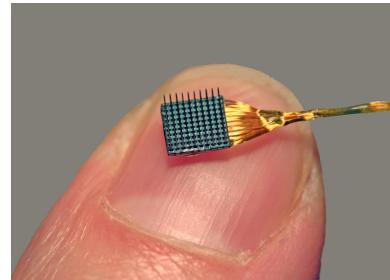
- ❑ Deal with large number of sensors/recording sites



- ❑ investigate high-dimensional representations
  - ❑ classification (what does this high-dimensional data represent?)
  - ❑ regression (how does it represent it? can we predict a different representation?)
  - ❑ model selection (what model would best describe this high dimensional data?)

# How can ML help neuroscientists?

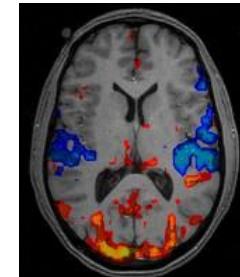
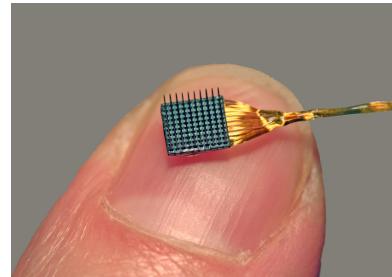
- ❑ Deal with large number of sensors/recording sites



- ❑ investigate high-dimensional representations
  - ❑ classification (what does this high-dimensional data represent?)
  - ❑ regression (how does it represent it? can we predict a different representation?)
  - ❑ model selection (what model would best describe this high dimensional data?)
- ❑ uncover few underlying processes that interact in complex ways

# How can ML help neuroscientists?

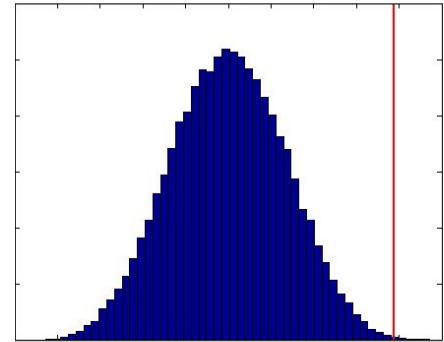
- ❑ Deal with large number of sensors/recording sites



- ❑ investigate high-dimensional representations
  - ❑ classification (what does this high-dimensional data represent?)
  - ❑ regression (how does it represent it? can we predict a different representation?)
  - ❑ model selection (what model would best describe this high dimensional data?)
- ❑ uncover few underlying processes that interact in complex ways
  - ❑ dimensionality reduction techniques

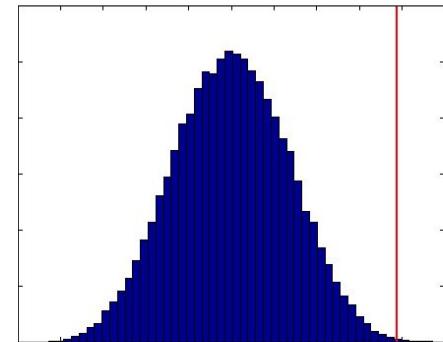
# How can ML help neuroscientists?

- ❑ Evaluate results



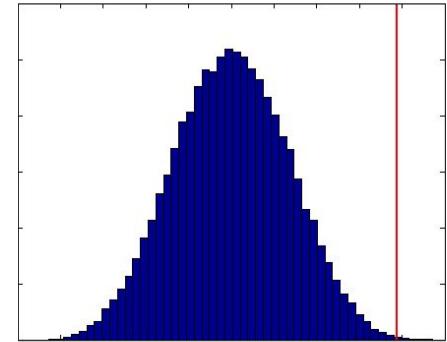
# How can ML help neuroscientists?

- ❑ Evaluate results
  - ❑ cross validation (how generalizable are our results?)



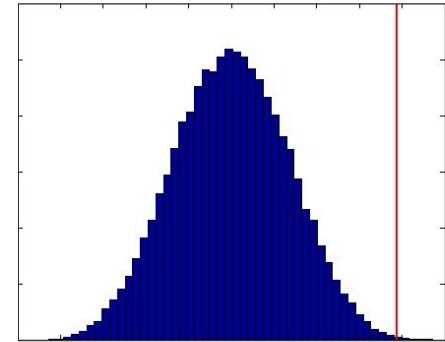
# How can ML help neuroscientists?

- ❑ Evaluate results
  - ❑ cross validation (how generalizable are our results?)
  - ❑ nearly assumption-free significance testing (are the results different from chance?)

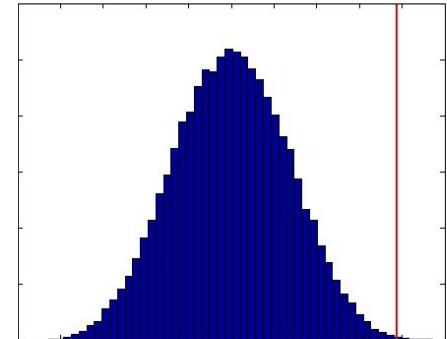


# How can ML help neuroscientists?

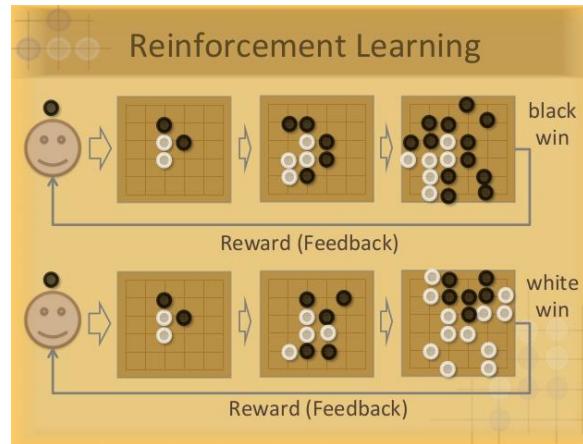
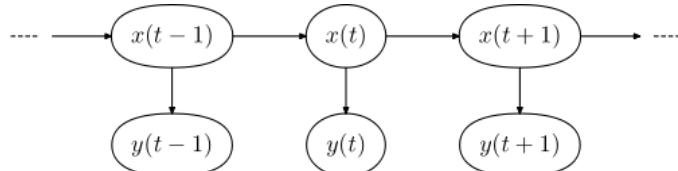
- ❑ Evaluate results
  - ❑ cross validation (how generalizable are our results?)
  - ❑ nearly assumption-free significance testing (are the results different from chance?)
- ❑ Complex data-driven hypotheses of brain processing



# How can ML help neuroscientists?



- ❑ Evaluate results
  - ❑ cross validation (how generalizable are our results?)
  - ❑ nearly assumption-free significance testing (are the results different from chance?)
- ❑ Complex data-driven hypotheses of brain processing
  - ❑ advanced topics: latent variable models, reinforcement learning, deep learning



# Thank you for your attention!

Please fill out the anonymous evaluation sheet -- your feedback is very important!

Looking forward to questions or follow-ups:

[mariya@cmu.edu](mailto:mariya@cmu.edu)